

Swarm Formal Specification

Swarm Research Division

draft of August 12, 2025

Contents

1	Formal specification	13
1.1	DISC basics	13
1.1.1	Notation	13
1.1.2	Sequences	13
1.1.3	Custom types	15
1.1.4	XOR distance and proximity order	16
1.1.5	Binary Merkle tree hash	17
1.1.6	Chunks	18
1.1.7	Single Owner Chunks	19
1.1.8	Segment inclusion proofs	20
1.1.9	Postage stamps	23
1.1.10	Ordering and sampling	23
1.2	Redistribution game	25
1.2.1	Transactions and on-chain registers	25
1.2.2	Random nonces	27
1.2.3	Winner selection and claim validation	28
1.2.4	Proofs of reserve	32
2	Data types and algorithms	37
2.1	Built-in primitives	37
2.1.1	Crypto	37
2.1.2	State store	43
2.1.3	Local context and configuration	44
2.2	Bzz address	44
2.2.1	Overlay address	44
2.2.2	Underlay address	45
2.2.3	BZZ address	45
2.3	Chunks, encryption and addressing	47
2.3.1	Content addressed chunks	47
2.3.2	Single owner chunk	50
2.3.3	Binary Merkle Tree Hash	51
2.3.4	Encryption	52

2.4	Postage stamps	54
2.4.1	Witness type	54
2.4.2	Contract Upgrades	55
2.5	Files, manifests and data structures	56
2.5.1	Files and the Swarm hash	56
2.5.2	Manifests	59
2.5.3	Resolver	63
2.5.4	Pinning	64
2.5.5	Tags	64
2.5.6	Storage	65
2.6	Access Control	66
2.7	PSS	69
2.7.1	PSS message	69
2.7.2	Direct pss message with trojan chunk	69
2.7.3	Envelopes	73
2.7.4	Update notifications	73
2.7.5	Chunk recovery	73
I	Appendix	75
A	Density-based size estimation	77
B	Source of randomness	83
C	Parameter constants	85

List of definitions

1	Definition (sequences)	13
2	Definition (segmentation)	14
3	Definition (Swarm overlay address of node n)	15
4	Definition (DISC custom types)	15
5	Definition (XOR distance (χ))	16
6	Definition (Proximity order (PO))	16
7	Definition (Binary Merkle Tree Root)	17
8	Definition (Binary Merkle Tree Hash)	18
9	Definition (Content addressed chunks)	18
10	Definition (Single owner chunks)	19
11	Definition (Packed address chunk)	20
12	Definition (BMT segment inclusion proof)	20
13	Definition (BMT prover function)	21
14	Definition (BMT SIP validation)	22
15	Definition (Single owner chunks data integrity proof)	22
16	Definition (Single owner chunks data integrity validation)	22
17	Definition (Postage stamps)	23
18	Definition (Storage slot reference)	23
19	Definition (Postage stamp validity)	23
20	Lemma (Ordering and indexing functions)	23
21	Corollary (hash orders)	24
22	Definition (Sampler function)	24
23	Definition (Redistribution game round)	25
24	Definition (Stakes)	25
25	Definition (Commits)	26
26	Definition (Reveals)	26
27	Definition (Random nonces for the round)	27
28	Definition (Prefixed hash)	28
29	Definition (Transformed chunk reserve sample)	29
30	Definition (Chunk reserve sample commitment hash)	29
31	Definition (Transformed slots reserve sample)	30
32	Definition (Transformed slots reserve sample commitment hash)	30

33	Definition (Weighted selection)	30
34	Definition (Truth selection)	31
35	Definition (Honest reveals)	31
36	Definition (Winner selection)	31
37	Definition (Proof of reserve)	32
38	Definition (Proof of reserve validation)	32
39	Definition (Proof of chunk density validation)	33
40	Definition (Proof of stamp density validation)	33
41	Definition (Proof of entitlement)	34
42	Definition (Winner's claim validation)	34
43	Corollary (Outpayment scheme is fair)	35
44	Definition (Hashing)	37
45	Definition (Random number generation)	38
46	Definition (Scripkey key derivation)	38
47	Definition (Mining a nonce)	39
48	Definition (Blockcipher)	39
49	Definition (Elliptic curve key generation)	40
50	Definition (Asymmetric encryption)	41
51	Definition (Signature)	41
52	Definition (Shared secret)	42
53	Definition (CRS erasure code interface definition)	43
54	Definition (CRS erasure coding parameters)	43
55	Definition (State store)	43
56	Definition (Context)	44
57	Definition (Swarm overlay address of node A)	44
58	Definition (Swarm underlay multiaddress)	45
59	Definition (Swarm bzz address transfer format)	45
60	Definition (Signed underlay address of node A)	46
61	Definition (Node addresses: overlay, underlay, bzz address)	46
62	Definition (segment, payload, span, branches)	47
63	Definition (Chunk reference)	48
64	Definition (Content addressed chunk)	48
65	Definition (Span to payload length)	49
66	Definition (Chunk API retrieval)	49
67	Definition (Chunk API: storage)	49
68	Definition (Single owner chunks)	50
69	Definition (Single owner chunk API: retrieval)	50
70	Definition (Single owner chunk API: storage)	51
71	Definition (BMT hash)	52
72	Definition (Chunk encryption/decryption API)	53
73	Definition (Postage contract)	55

74	Definition (Postage stamp basic types: batchID, address, witness, stamp, validity)	55
75	Definition (File API: upload/storage; swarm hash)	56
76	Definition (File API: download/retrieval)	58
77	Definition (File info)	59
78	Definition (Manifest entry)	59
79	Definition (Manifest data structure)	60
80	Definition (Manifest API: path lookup)	60
81	Definition (Manifest API: update)	61
82	Definition (Manifest API: remove)	61
83	Definition (Manifest API: merge)	62
84	Definition (Resolver)	63
85	Definition (Pinning)	64
86	Definition (Tags)	64
87	Definition (Public storage API)	65
88	Definition (Access control: auth, hint, parameters, root access manifest) . . .	66
89	Definition (Session key and auth with credentials)	67
90	Definition (Access key)	67
91	Definition (Access control API: lock/unlock)	68
92	Definition (ACT manipulation API: add/remove)	68
93	Definition (Basic types: topic, targets, recipient, message and trojan)	69
94	Definition (Sealing/unsealing the message)	70
95	Definition (Wrapping/unwrapping)	71
96	Definition (Incoming message handling)	71
97	Definition (pss API: send)	72
98	Definition (pss API: receive)	72
99	Definition (Envelope)	73
100	Definition (Targeted delivery request)	73
101	Definition (Prod missing chunk notification with recovery request)	74
102	Definition (Recovery response)	74
103	Definition (Random seed for the round)	84
104	Lemma (Round seed is a secure random oracle)	84

List of Figures

1.1	BMT: Binary Merkle Tree hash	17
1.2	Compact segment inclusion proofs for chunks	21
A.1	Probability density function of the Erlang distribution	79
A.2	The 16th smallest value among 500 random variates on unit interval	79
A.3	Recall and precision of proof of reserve size validation	81
A.4	Optimal accuracy of reserve size probe	82

List of Tables

2.1	Size of chunk spans	56
A.1	Proof of density parameter calibration	82
C.1	Parameter constants	86

Chapter 1

Formal specification

1.1 DISC basics

1.1.1 Notation

- roman – predicates standing for blockchain verifiable properties
- SMALLCAPS – field of a composite, tuple member function
- Π^p – proof constructor function
- \mathbb{V}^p – validator function
- `TT_FONT` – constant parameter (see appendix ??)
- $\text{uint}[d]$ – left closed, right open integer range $\overline{[0, 2^d)}$

1.1.2 Sequences

Definition 1 – sequences .

Define $\tau\{n\}$, the non-polymorphic sequences of length n (non-negative) over any type τ as an indexing function:

$$\tau\{n\} \stackrel{\text{def}}{=} \begin{cases} \emptyset & \text{if } n = 0 \\ \overline{0, n-1} \rightarrow \tau & \text{if } n > 0 \end{cases} \quad (1.1)$$

$$\tau+ \stackrel{\text{def}}{=} \bigcup_{n \in \mathbb{Z}^+} \tau\{n\} \quad (1.2)$$

$$\tau* \stackrel{\text{def}}{=} \{\emptyset\} \cup \tau+ \quad (1.3)$$

$$(1.4)$$

Length function len :

$$len(s) : \tau \rightarrow uint64 \quad (1.5)$$

$$len(s) \stackrel{\text{def}}{=} \begin{cases} 0 & \text{if } s = \emptyset \\ n & \text{if } s \in \tau\{n\} \end{cases} \quad (1.6)$$

The positional index 'at' function:

$$[] : \tau^* \rightarrow \times uint64 \rightarrow \tau \quad (1.7)$$

$$s[] : \overline{0, len(s) - 1} \rightarrow \tau \quad (1.8)$$

$$s[i] \stackrel{\text{def}}{=} s'(i) \quad (1.9)$$

Define the *concatenation* operator ' \oplus ':

$$(x \oplus y) : \overline{0, len(x) + len(y) - 1} \rightarrow \tau \quad (1.10)$$

$$(x \oplus y)[i] \stackrel{\text{def}}{=} \begin{cases} x[i] & \text{if } 0 \leq i < len(x) \\ y[i - len(x)] & \text{otherwise} \end{cases} \quad (1.11)$$

Define slices (subsequences) with the range operator ':':

$$(s[o : o + l]) : \overline{0, l - 1} \rightarrow \tau \quad (1.12)$$

$$(s[o : o + l])[i] \stackrel{\text{def}}{=} s[o + i] \quad (1.13)$$

where

$$o, l \geq 0 \wedge \quad (1.14)$$

$$len(x) \geq o + l \quad (1.15)$$

Let us also define empty, prefix and suffix slices:

$$s[x : x] \stackrel{\text{def}}{=} \emptyset \quad (1.16)$$

$$s[:x] \stackrel{\text{def}}{=} s[0 : x] \quad (1.17)$$

$$s[x:] \stackrel{\text{def}}{=} s[x : len(s)] \quad (1.18)$$

As a special case byte slices are defined as sequences of 8-bit integers.

Definition 2 – segmentation.

Define segment as an at most 32 long byte slice, and define segmentation of a slice of bytes as the partitioning of the slice into consecutive segments. Define segment count as the number of segments that cover a byte slice:

$$SegCnt : byte^* \rightarrow uint64 \quad (1.19)$$

$$SegCnt(s) \stackrel{\text{def}}{=} int \left(\frac{len(s) - 1}{32} \right) + 1 \quad (1.20)$$

Now we can define the segment indexing function $[[[]]]$ which maps the byte slice s and an index i to the i -th segment in the segmentation of s :

$$[[[]]] : \text{byte}^* \rightarrow \text{uint64} \rightarrow \text{Segment} \quad (1.21)$$

$$s[[[]]] : \overline{0, \text{SegCnt}(s) - 1} \rightarrow \text{Segment} \quad (1.22)$$

$$s[[i]] \stackrel{\text{def}}{=} \begin{cases} s[32 \cdot i:] & \text{if } i = \text{SegCnt} - 1 \\ s[32 \cdot i : 32 \cdot (i + 1)] & \text{otherwise} \end{cases} \quad (1.23)$$

1.1.3 Custom types

Definition 3 – Swarm overlay address of node n .

A Swarm node is associated with an Ethereum account that the node operator must possess the private key for, called *bzz account* (K_n^{bzz}). The node's overlay address is derived as the hash of the binary serialisation of the Ethereum address of this account with the Swarm network ID and a minable nonce appended.

$$\text{overlay}(n) \stackrel{\text{def}}{=} H(\text{acc} \oplus \text{id} \oplus \nu) \quad (1.24)$$

where

$$\text{ETH ADDRESS } \text{acc} = \text{Account}(K_n) \quad (1.25)$$

$$\text{NETWORK ID } \text{id} = \text{BZZ_NETWORK_ID} \quad (1.26)$$

$$\text{OVERLAY NONCE } \nu \in \text{Nonce} \quad (1.27)$$

Definition 4 – DISC custom types .

Let us define the DISC specific custom types used in the formalisation.

$\text{Segment} \equiv \text{byte}\{32\} \equiv \text{uint256}$

32-long slice of raw bytes

in numerical context cast as BigEndian encoded 256-bit unsigned integer

$\text{Address} \equiv \text{Segment}$

swarm chunk address, swarm peers' overlay address

$\text{Nonce} \equiv \text{Segment}$

deterministically random Segment

$\text{Account} \equiv \text{byte}\{20\}$

Ethereum address derived from EC keypair K

$\text{Account}(K) \stackrel{\text{def}}{=} H(\text{PubKey}(K))[12:32]$

$\text{Nodes} \equiv \text{Segment}$

swarm client node (peer)

$Sig \equiv \text{byte}\{65\}$

$\langle r, s, v \rangle$ representation of an EC signature (32+32+1 bytes)

$Timestamp \equiv \text{uint64}$

64-bit unsigned integer for unix time, nanosecond resolution

big endian binary serialisation.

$H : \text{byte}^* \rightarrow \text{Segment}$

the 256-bit Keccak SHA3 hash function, the base hash used in swarm.

1.1.4 XOR distance and proximity order

Definition 5 – XOR distance (χ).

Consider the set of bit sequences with fixed length d as points in a space. Define a distance metric χ such that the distance between two such sequences is the numerical value of their bitwise XOR (\wedge) using big endian (= most significant bit first) encoding.

$$\chi : \text{uint}[d] \times \text{uint}[d] \rightarrow \text{uint}[d] \quad (1.28)$$

$$\chi(x, y) \stackrel{\text{def}}{=} \text{Uint}^d(\text{BE}^d(x) \wedge \text{BE}^d(y)) \quad (1.29)$$

Given the fixed length $d > 0$, there is a maximum distance ($2^d - 1 = \chi(0\{d\}, 1\{d\})$) in this space, and thus we can define the notion of *normalised distance*:

$$\bar{\chi} : \text{uint}[d] \times \text{uint}[d] \rightarrow \mathbb{Q}[0, 1] \quad (1.30)$$

$$\bar{\chi}(x, y) \stackrel{\text{def}}{=} \frac{\chi(x, y)}{2^d - 1} \quad (1.31)$$

Definition 6 – Proximity order (PO).

Proximity order (PO) is a discrete logarithmic scaling of proximity.

$$PO : \text{uint}[d] \times \text{uint}[d] \rightarrow \overline{0, d} \quad (1.32)$$

$$PO(x, y) \stackrel{\text{def}}{=} \begin{cases} d & \text{if } x = y \\ \text{int}(\log_2(\text{Proximity}(x, y))) & \text{otherwise} \end{cases} \quad (1.33)$$

where proximity is the inverse of normalised distance:

$$\text{Proximity} : \text{uint}[d] \times \text{uint}[d] \rightarrow \text{uint}[d] \quad (1.34)$$

$$\text{Proximity}(x, y) \stackrel{\text{def}}{=} \frac{1}{\bar{\chi}(x, y)} \quad (1.35)$$

Given two points x and y , the order of their proximity $PO(x, y)$ equals the number of initial bits shared by their respective most significant bit first binary representations. In practice, with $d = 256$, $uint[d] \equiv Segment$, so PO also applies to a pair of slices of 32 bytes.

1.1.5 Binary Merkle tree hash

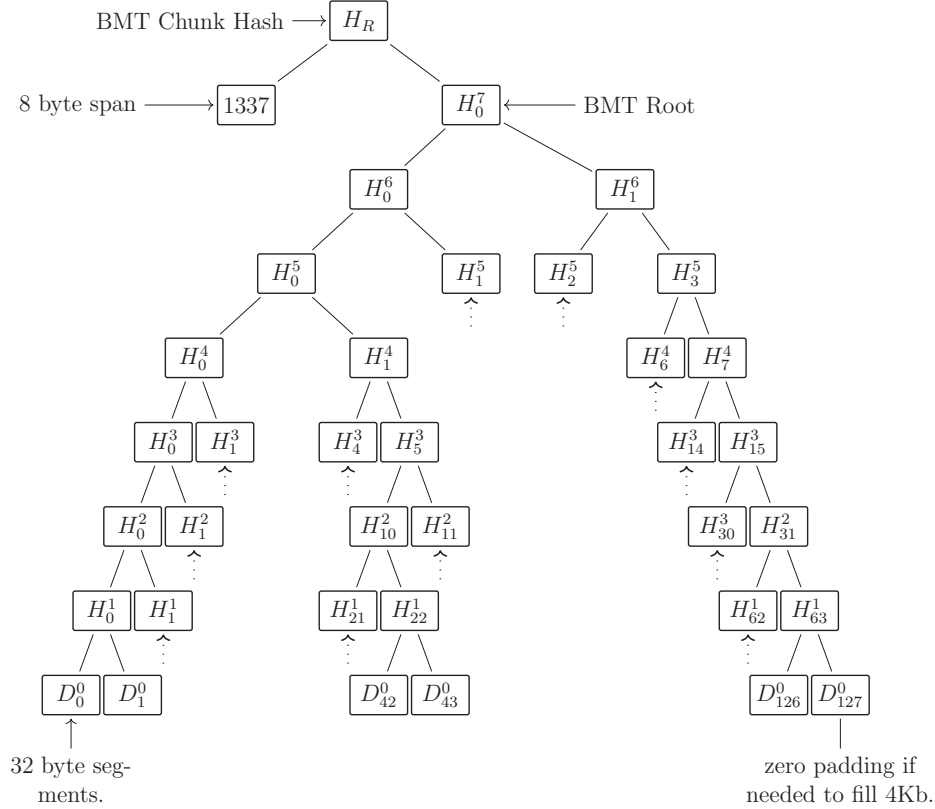


Figure 1.1: BMT (Binary Merkle Tree) used as the native chunk hash in Swarm. In this example, 1337 bytes of chunk data is segmented into 32 byte segments. Zero padding is used to fill up the rest up to 4 kilobytes. Pairs of segments are hashed together using the ethereum-native Keccak256 hash to build up the binary tree. On level 8, the binary Merkle root is prepended with the 8 byte span and hashed to yield the BMT chunk hash.

The BMT chunk address is the hash of the 8 byte metadata (span) and the root hash of a *binary Merkle tree* (BMT) built on the 32-byte segments of the underlying data (see figure 1.1). If the chunk content is less than 4k, the hash is calculated as if the chunk was padded with all zeros up to 4096 bytes.

Definition 7 – Binary Merkle Tree Root .

Define $\Delta[H, n](i, d)$ as the Binary Merkle Tree Root of data d fitting in at most 2^n 32-byte

segments using H as its base hash:

$$\Delta : (\text{byte}^* \rightarrow \text{Segment}) \times \text{uint8} \rightarrow \text{uint8} \times \text{byte}^* \rightarrow \text{Segment} \quad (1.36)$$

$$\Delta[H, n] = : \overline{0, n} \times \text{byte}\{, 2^n\} \rightarrow \text{Segment} \quad (1.37)$$

$$\Delta[H, n](i, d) \stackrel{\text{def}}{=} \begin{cases} d & \text{if } i = 0 \\ \Delta[H, n](i, d \oplus 0\{2^{n+5} - \text{len}(d)\}) & \text{if } \text{len}(d) < 2^{i+5} \\ H(\Delta[H, n](d[:2^{i-1}], i-1) \oplus \\ \oplus \Delta[H, n](d[2^{i-1}:], i-1)) & \text{otherwise} \end{cases} \quad (1.38)$$

Definition 8 – Binary Merkle Tree Hash .

Define $BMT[H](d, m)$ as the hash of Binary Merkle Tree Root of chunk length data d prepended with metadata m . H is the base hash function, by default 256-bit Keccak. Note that a chunk size blob of bytes can accommodate 128 32-byte segments, hence depth 7:

$$BMT : (\text{byte}^* \rightarrow \text{Segment}) \rightarrow \text{Chunk} \times \text{byte}\{8\} \rightarrow \text{Segment} \quad (1.39)$$

$$BMT[H] : \text{Chunk} \times \text{byte}\{8\} \rightarrow \text{Segment} \quad (1.40)$$

$$BMT[H](d, m) \stackrel{\text{def}}{=} H(m \oplus \Delta[H, 7](7, d)) \quad (1.41)$$

1.1.6 Chunks

Definition 9 – Content addressed chunks .

Define content address chunk c as a function from bytes data with size limit of 4096 bytes and an associated address calculated with BMT:

$$CAC : \text{Chunk} \times \text{byte}\{8\} \rightarrow \text{Chunks} \quad (1.42)$$

$$CAC(d, m) \stackrel{\text{def}}{=} \langle \text{addr}, \text{cont} \rangle \quad (1.43)$$

such that

$$\text{addr} \stackrel{\text{def}}{=} BMT(d, m) \quad (1.44)$$

$$\text{cont} \stackrel{\text{def}}{=} m \oplus d \quad (1.45)$$

By convention the metadata prefix m encodes the *span* using 64-bit little endian. If the chunk is an intermediate chunk (see definition 11), the span is the length of the data that the subtree spans over. If the chunk is a data chunk, then span encodes the data length:

$$m \stackrel{\text{def}}{=} LE^{64}(\text{len}(d)) \quad (1.46)$$

We also say that for $cac = CAC(d, m) = \langle addr, cont \rangle$:

$$\text{ADDRESS}(cac) \stackrel{\text{def}}{=} addr \quad (1.47)$$

$$\text{PAYLOAD}(cac) \stackrel{\text{def}}{=} cont \quad (1.48)$$

$$\text{DATA}(cac) \stackrel{\text{def}}{=} d \quad (1.49)$$

$$\text{METADATA}(cac) \stackrel{\text{def}}{=} m \quad (1.50)$$

For convenience $SegCnt$ and the segment indexing function ' $[[\cdot]]$ ' can be trivially extended to apply to chunks:

$$SegCnt : Chunks \rightarrow uint64 \quad (1.51)$$

$$SegCnt(c) \stackrel{\text{def}}{=} SegCnt(\text{DATA}(c)) \quad (1.52)$$

$$[[\cdot]] : Chunks \times uint64 \rightarrow Segment \quad (1.53)$$

$$c[[i]] \stackrel{\text{def}}{=} \text{DATA}(c)[[i]] \quad (1.54)$$

1.1.7 Single Owner Chunks

Definition 10 – Single owner chunks .

A single owner chunk is defined as a content addressed chunk associated with an ID and an ethereum address:

$$SOC \stackrel{\text{def}}{=} Account \times Segment \times CAC \rightarrow Chunks \quad (1.55)$$

$$SOC(owner, id, cac) \stackrel{\text{def}}{=} \langle addr, cont \rangle \quad (1.56)$$

where

$$addr \stackrel{\text{def}}{=} H(o \oplus id) \quad (1.57)$$

$$cont \stackrel{\text{def}}{=} id \oplus Sig(o, id \oplus \text{ADDRESS}(cac)) \oplus \text{PAYLOAD}(cac) \quad (1.58)$$

A single owner chunk's address is the Keccak256 hash of identifier prepended to owner account, while its data is serialised as follows:

- *identifier* – 32 bytes arbitrary identifier,
- *signature* – 65 bytes $\langle r, s, v \rangle$ representation of an EC signature (32+32+1 bytes),

- *span* – 8 byte little endian binary of uint64 chunk span,
- *data* – max 4096 bytes of regular chunk data.

Integrity of a *single owner chunk* is verified with the following process:

1. Deserialise the chunk content into fields for identifier, signature and payload.
2. Construct the expected plaintext composed of the identifier and the *BMT hash* of the payload.
3. Recover the owner's address from the signature using the plaintext.
4. Check the hash of the identifier and the owner (expected address) against the chunk address.

Definition 11 – Packed address chunk.

Define the packed address chunk for a sequence of chunks C as the concatenation of all the addresses of the chunks in the sequence:

$$PAC : Chunks* \times \text{byte}\{8\} \rightarrow Chunks \quad (1.59)$$

$$PAC(C, m) \stackrel{\text{def}}{=} CAC \left(\bigoplus_{i=0}^{\text{len}(C)-1} \text{Address}(C[i]), m \right) \quad (1.60)$$

1.1.8 Segment inclusion proofs

Using BMT hashes allows for compact *segment inclusion proofs* (substring relationship with a 32-byte resolution).

Definition 12 – BMT segment inclusion proof .

Define $\Pi^{\text{SIP}}(c, i)$ as the *BMT* inclusion proof on chunk c for segment index i :

$$\Pi^{\text{SIP}} : Chunks \times \overline{0, 127} \rightarrow SIP \quad (1.61)$$

$$SIP \stackrel{\text{def}}{=} Segment \times Segment^7 \times \text{byte}\{8\} \quad (1.62)$$

$$\Pi^{\text{SIP}}(c, i) \stackrel{\text{def}}{=} \langle c[[i]], \langle h_0, h_1, \dots, h_6 \rangle, \text{METADATA}(c) \rangle \quad (1.63)$$

where

$$h_j \stackrel{\text{def}}{=} BMT(s_j, j) \quad (1.64)$$

$$s_j \stackrel{\text{def}}{=} c[\text{start}(i, j) : \text{start}(i, j) + 32 \cdot 2^j] \quad (1.65)$$

where

$$\text{start}(i, j) \stackrel{\text{def}}{=} \begin{cases} 0 & \text{if } j = 7 \\ \text{start}(i, j+1) & \text{if } \text{int}(i/2^j) = 0 \pmod{2} \\ \text{start}(i, j+1) + 32 \cdot 2^j & \text{otherwise} \end{cases} \quad (1.66)$$

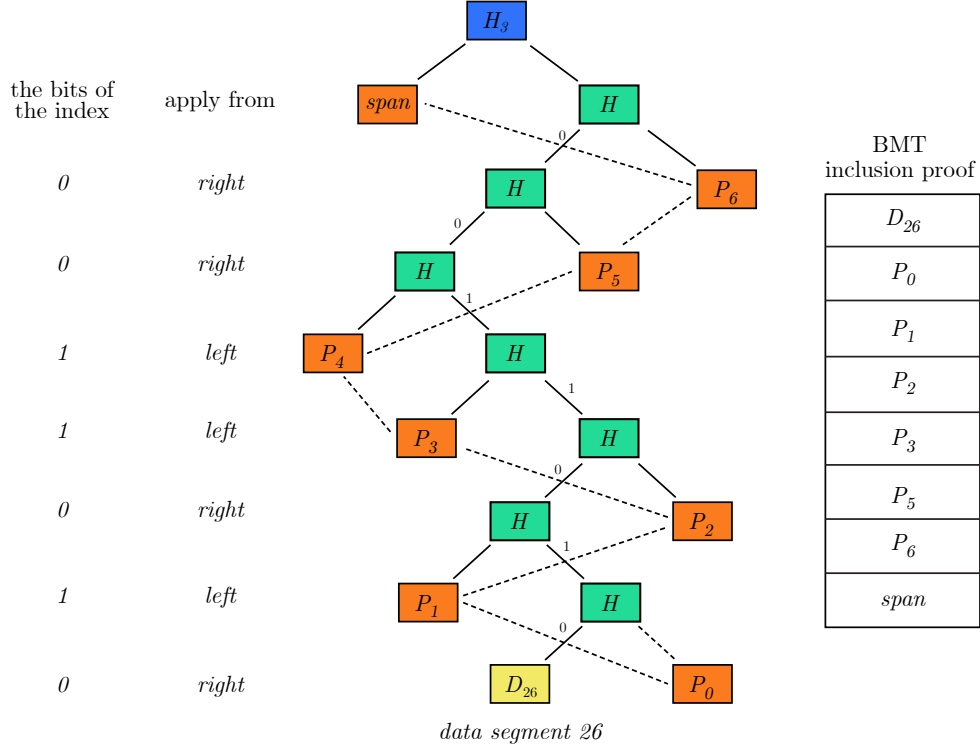


Figure 1.2: Compact segment inclusion proofs for chunks. Assume we need proof for segment 26 of a chunk (yellow). The orange hashes of the BMT are the sister nodes on the path from the data segment up to the root and constitute what needs to be part of a proof. When these are provided together with the root hash and the segment index, the proof can be verified. The side on which proof item i needs to be applied depends on the i -th bit (starting from least significant) of the binary representation of the index. Finally the span is prepended and the resulting hash should match the chunk root hash.

In order to validate segment inclusion proofs we first introduce the prover hash function H_{Π} .

Definition 13 – BMT prover function.

$$H_{\Pi} : \overline{0, 127} \times SIP \rightarrow Address \quad (1.67)$$

$$H_{\Pi}(i, \langle d, sisters, m \rangle) \stackrel{\text{def}}{=} H(m, H_{\Pi}^{\Delta}(7, d, sisters)) \quad (1.68)$$

where

$$H_{\Pi}^{\Delta} : \overline{0,7} \times \text{Segment} \times \text{Segment}^7 \rightarrow \text{Address} \quad (1.69)$$

$$H_{\Pi}^{\Delta}(j, d, s) \stackrel{\text{def}}{=} \begin{cases} d & \text{if } j = 0 \\ H(H_{\Pi}^{\Delta}(j-1, d, s) \oplus s[j-1]) & \text{if } \text{int}(i/2^{j-1}) = 0 \pmod{2} \\ H(s[j-1] \oplus H_{\Pi}^{\Delta}(j-1, d, s)) & \text{otherwise} \end{cases} \quad (1.70)$$

Definition 14 – BMT SIP validation .

Define $\mathbb{V}^{\text{SIP}}(a, i, p)$ as the validator of a BMT segment inclusion proof p for chunk at address a on segment index i :

$$\mathbb{V}^{\text{SIP}} : \text{Segment} \times \overline{0,127} \times \text{SIP} \rightarrow \{\text{T}, \text{F}\} \quad (1.71)$$

$$\mathbb{V}^{\text{SIP}}(a, i, p) \Leftrightarrow H_{\Pi}(i, p) = a \quad (1.72)$$

Definition 15 – Single owner chunks data integrity proof .

Define a single owner chunk storage proof $\Pi^{\text{SOC}}(c, i)$ as a segment inclusion proof of the data payload of SOC c on index i together with the ID and signature of SOC c :

$$\text{SIP}_{\text{SOC}} \stackrel{\text{def}}{=} \text{SIP} \times \text{Sig} \times \text{Segment} \quad (1.73)$$

$$\Pi^{\text{SIP}[\text{SOC}]} : \text{SOC} \times \overline{0,127} \rightarrow \text{SIP}_{\text{SOC}} \quad (1.74)$$

$$\Pi^{\text{SIP}[\text{SOC}]}(\langle o, id, cac \rangle, i) \stackrel{\text{def}}{=} \langle p, sig, id \rangle \quad (1.75)$$

where

$$p = \Pi^{\text{SIP}}(cac, i) \quad (1.76)$$

$$sig = \text{Sig}(o, id \oplus \text{address}(cac)) \quad (1.77)$$

Definition 16 – Single owner chunks data integrity validation .

Define $\mathbb{V}^{\text{SIP}[\text{SOC}]}(a, i, p)$ as the validator of a single owner chunk storage proof p for chunk at address a on segment index i :

$$\mathbb{V}^{\text{SIP}[\text{SOC}]} : \text{Address} \times \overline{0,127} \times \text{SIP}_{\text{SOC}} \rightarrow \{\text{T}, \text{F}\} \quad (1.78)$$

$$\mathbb{V}^{\text{SIP}[\text{SOC}]}(a, i, \langle p, sig, id \rangle) \Leftrightarrow a = H(id \oplus o) \quad (1.79)$$

such that

$$\text{OWNER } o = \text{ECRecover}(sig, id \oplus a') \quad (1.80)$$

$$\text{PAYLOAD } a' = H_{\Pi}(i, p) \quad (1.81)$$

1.1.9 Postage stamps

Definition 17 – Postage stamps .

$$Stamps \stackrel{\text{def}}{=} Segment \times uint64 \times Timestamp \times Address \quad (1.82)$$

$$ps = \langle b, i, ts, a \rangle \in Stamps \quad (1.83)$$

$$BATCHID(ps) \stackrel{\text{def}}{=} b \quad (1.84)$$

$$INDEX(ps) \stackrel{\text{def}}{=} i \quad (1.85)$$

$$TIMESTAMP(ps) \stackrel{\text{def}}{=} ts \quad (1.86)$$

$$ADDRESS(ps) \stackrel{\text{def}}{=} a \quad (1.87)$$

Definition 18 – Storage slot reference .

Define the *storage slot reference* $slot(ps)$ of a postage stamp ps as the tuple of the batch identifier and the within-batch stamp counter:

$$Slots \stackrel{\text{def}}{=} Segment \times uint64 \quad (1.88)$$

$$slot : Stamps \rightarrow Slots \quad (1.89)$$

$$slot(ps) \stackrel{\text{def}}{=} \langle BATCHID(ps), INDEX(ps) \rangle \quad (1.90)$$

Definition 19 – Postage stamp validity .

Define $\mathbb{V}^{\text{STAMP}}(ps)$ as the validator of the proof of relevance expressed as the postage stamp ps relying on blockchain information:

$$\mathbb{V}^{\text{STAMP}} : Stamps \times \Gamma \times Nodes \rightarrow \{T, F\} \quad (1.91)$$

$$\mathbb{V}^{\text{STAMP}}(ps, \gamma, n) \Leftrightarrow \quad (1.92)$$

$$\text{AUTHENTIC} \quad BATCHID(ps) \in \text{Batches}(\gamma) \wedge \quad (1.93)$$

$$\text{ALIVE} \quad \text{Balance}(ps) > 0 \wedge \quad (1.94)$$

$$\text{AUTHORISED} \quad ECR\text{recover}(Sig(ps), encode(ps)) = \text{Owner}(ps) \wedge \quad (1.95)$$

$$\text{AVAILABLE} \quad 0 \leq INDEX(ps) < \text{Size}(ps) \wedge \quad (1.96)$$

$$\text{ALIGNED} \quad PO(INDEX(ps), n) \geq \text{DEPTH}(reveal(\gamma, n)) \quad (1.97)$$

1.1.10 Ordering and sampling

Lemma 20 – Ordering and indexing functions.

Given an arbitrary finite set C , and another set I with a total order $<$. Any invertible

function total over C , $f : C \rightarrow I$ defines a total order $<_f$ over C as follows:

$$<_f \subseteq C \times C \quad (1.98)$$

$$\forall c, c' \in C, c <_f c' \Leftrightarrow f(c) < f(c') \quad (1.99)$$

Proof. f is injective wrt I , so $\text{Image}(f) = J \subseteq I$. Since $<$ restricted to a subset ($<_J$) is also a total order over J . Since f is invertible, C and J are isomorphic and therefore the total order $<$ on J carries over to C .

Corollary 21 – hash orders.

Any hash function defines a total order on a finite set of byte slices.

Proof. The collision free nature of the hash function makes it practically invertible. The actual hashes when read as binary encodings of integers, offer a natural integer ordering over the values.

Example: the prefixed BMT hash transform (see 28) defines a total order over a set of chunks.

Definition 22 – Sampler function.

Using f and its derivative ordering on $C \subseteq \text{Dom}(f)$ we represent C as an ordered sequence:

$$\overrightarrow{\text{Seq}} : (T \rightarrow T) \times \mathcal{P}(T) \rightarrow T^* \quad (1.100)$$

$$\overrightarrow{\text{Seq}}(f, C) \stackrel{\text{def}}{=} C' \quad (1.101)$$

$$\text{such that } f(C'[i]) < f(C'[j]) \quad \text{for every } 0 \leq i < j < |C| \quad (1.102)$$

Finally, we define a sampler function for any f invertible with an image having a total order and $C \subseteq \text{Dom}(f)$ such that it selects a prefix slice of length l from the ordered C :

$$\text{Sampler} : (T \rightarrow T) \times \mathcal{P}(T) \times \text{uint64} \rightarrow T^+ \quad (1.103)$$

$$\text{Sampler}(f, C, l) \stackrel{\text{def}}{=} \overrightarrow{\text{Seq}}(f, C)[: l] \quad (1.104)$$

1.2 Redistribution game

Definition 23 – Redistribution game round.

Define γ as a redistribution game round. γ is conceived of as a multidimensional index:

$$\Gamma \stackrel{\text{def}}{=} \text{uint64} \times \text{uint64} \times \text{uint64} \quad (1.105)$$

$$\gamma \in \Gamma = \langle c, \sigma, i \rangle \quad (1.106)$$

$$\text{CHAIN}(\gamma) \quad c \quad \text{ID of the blockchain context} \quad (1.107)$$

$$\text{SERIES}(\gamma) \quad \sigma \quad \text{index of the parallel series} \quad (1.108)$$

$$\text{ROUND}(\gamma) \quad i \quad \text{sequential index of the round} \quad (1.109)$$

Define $\text{BLOCK}(\gamma)$ as the starting block height of this particular game γ :

$$\text{BLOCK}(\gamma) \stackrel{\text{def}}{=} \text{ROUND}(\gamma) \cdot \text{ROUND_LENGTH} + \text{START_BLOCK} \quad (1.110)$$

Ordering by sequential index defines the chain of games, which lets us define the *Prev* function:

$$\text{Prev} \quad : \quad \Gamma \rightarrow \Gamma \quad (1.111)$$

$$\text{Prev}(\langle c, \sigma, i \rangle) \stackrel{\text{def}}{=} \langle c, \sigma, i - 1 \rangle \quad (1.112)$$

1.2.1 Transactions and on-chain registers

The smart contract receives transactions from applicants in phases. The following virtual registers capture the information given in these transactions that are relevant for defining the winner:

- batches (see definition 19)
- stakes (see definition 24)
- commits (see definition 25)
- reveals (see definition 26)

Definition 24 – Stakes .

We define *Stakes* as the registry of stakes resulting from transactions sent to the staking contract. A record is a tuple of a node overlay, the stake balance and the committed stake

and can be updated.

$$Stakes : \Gamma \rightarrow Nodes \times uint64 \times uint64 \quad (1.113)$$

$$\langle n, s, m \rangle \in Stakes(\gamma) \Leftrightarrow \quad (1.114)$$

$$RIGHT\ AGE \quad \exists b' < b - MIN_STAKE_AGE, \tau \in Transactions(b') \quad (1.115)$$

$$NODE\ OVERLAY \quad n = H(origin(\tau) \oplus BZZ_NETWORK_ID \oplus data(\tau)[0]) \quad (1.116)$$

$$STAKE\ BALANCE \quad s = amount(\tau) \quad (1.117)$$

$$COMMITTED\ STAKE \quad m = data(\tau)[1] \quad (1.118)$$

There is only one stake allowed per node, so we can define the staked amount belonging to a node as the minimum of the stake balance and the committed stake times the unit price of storage:

$$Stake : \Gamma \times Nodes \rightarrow uint64 \quad (1.119)$$

$$Stake(\gamma, n) \stackrel{\text{def}}{=} min(s, m \cdot Price(\gamma)) \quad (1.120)$$

Definition 25 – Commits .

We define $Commits(\gamma)$ as the registry of applications for a game γ resulting from a transaction sent to the game contract's *commit* endpoint. A tuple of the overlay of the committing node, its commitment hash and the number of the block containing the transaction is entered in the register after verifying that (i) the transaction was sent during the commit phase (right time), and (ii) that the node has enough stake and is not frozen (right amount).

$$Commits : \Gamma \rightarrow Nodes \times Segment \times Blocks \quad (1.121)$$

$$\langle n, h, b \rangle \in Commits(\gamma) \Leftrightarrow \quad (1.122)$$

$$RIGHT\ TIME \quad b < PHASE_LENGTH \mod ROUND_LENGTH \quad (1.123)$$

$$RIGHT\ AMOUNT \quad stake(\gamma, n) \geq MINIMUM_STAKE \quad (1.124)$$

$$(REDUNDANCY) \quad (1.125)$$

Definition 26 – Reveals.

We define *Reveals* as the registry of reveals resulting from a transaction sent to the game contract's reveal endpoint. The reveal record is a tuple of the node overlay, the two commitment hashes, the self-reported storage depth, a serial index used for sorting, the obfuscation key, and the block number. The record is entered in the register after it is validated that (i) it was submitted during the reveal phase (right time), (ii) the commitments when obfuscated match the commit by the same node (right reveal), and (iii) that the neighbourhood selection anchor falls within the node's area of responsibility using the self-reported depth

(right location).

$$RevealEntry : Nodes \times Address^2 \times uint8^2 \times Nonce \times Block \quad (1.126)$$

$$Reveals : \Gamma \rightarrow \mathcal{P}(RevealEntry) \quad (1.127)$$

$$r = \langle n, chc, chs, sd, i, k, b \rangle \in Reveals(\gamma) \Leftrightarrow \langle \quad (1.128)$$

$$NODE(r) = n \quad (1.129)$$

$$CHC(r) = chc \quad (1.130)$$

$$CHS(r) = chs \quad (1.131)$$

$$DEPTH(r) = sd \quad (1.132)$$

$$INDEX(r) = i \quad (1.133)$$

$$NONCE(r) = k \quad (1.134)$$

$$BLOCK(r) = b \quad (1.135)$$

$$\rangle \quad (1.136)$$

$$\Leftrightarrow \quad (1.137)$$

$$\text{RIGHT TIME} \quad p \leq b < 2p \mod r \quad (1.138)$$

$$p = \text{PHASE_LENGTH}, r = \text{ROUND_LENGTH} \quad (1.139)$$

$$\text{RIGHT REVEAL} \quad H(n \oplus sd \oplus chc \oplus chs \oplus k) = h \text{ such that } (1.140)$$

$$\langle n, h, b' \rangle \in Commits(\gamma) \text{ for some } b' \quad (1.141)$$

$$\text{RIGHT LOCATION} \quad PO(addr(n), NSA(Prev(\gamma))) \geq sd \quad (1.142)$$

$$(\text{RESPONSIBILITY}) \quad (1.143)$$

There is only one reveal allowed per node, so we can define the reveal belonging to a node:

$$Reveal : \Gamma \times Nodes \rightarrow RevealEntry \quad (1.144)$$

$$Reveal(\gamma, n) = r \quad (1.145)$$

such that

$$n = Node(r) \quad (1.146)$$

$$r \in Reveals(\gamma) \quad (1.147)$$

1.2.2 Random nonces

Definition 27 – Random nonces for the round.

From the round's random seed (see definition 103 appendix ??) we can derive all the necessary

random input nonces:

$$\text{N.HOOD SELECTION ANCHOR } NSA(\gamma) \stackrel{\text{def}}{=} H(\mathcal{R}(\gamma) \oplus BE^{64}(SK(\gamma))) \quad (1.148)$$

$$\text{TRUTH SELECTION NONCES } TSN(\gamma) \stackrel{\text{def}}{=} H(\mathcal{R}(\gamma) \oplus BE^8(0)) \quad (1.149)$$

$$\text{WINNER SELECTION NONCES } WSN(\gamma, i) \stackrel{\text{def}}{=} H(\mathcal{R}(\gamma) \oplus BE^8(1)) \quad (1.150)$$

$$\text{RESERVE SAMPLING SALT } RSS(\gamma) \stackrel{\text{def}}{=} H(\mathcal{R}(\gamma)) \quad (1.151)$$

$$\text{SEGMENT SELECTION NONCE } SSN(\gamma, i) \stackrel{\text{def}}{=} H(\mathcal{R}(\gamma) \oplus BE^8(i)) \quad (1.152)$$

where

$$SK : \Gamma \rightarrow uint64 \quad (1.153)$$

$$SK(\gamma) \stackrel{\text{def}}{=} \begin{cases} 0 & \text{if } |Reveals(\gamma)| > 0 \\ SK(Prev(\gamma)) + 1 & \text{otherwise} \end{cases} \quad (1.154)$$

Let us now define the witness selection function \mathcal{W} that selects two random witness indexes as well as the last index of the reserve sample such that they are all distinct:

$$\mathcal{W} : \Gamma \times uint64 \times \{0, 1, 2\} \rightarrow uint8 \quad (1.155)$$

$$\mathcal{W}(\gamma, m, k) \stackrel{\text{def}}{=} \begin{cases} SSN(\gamma, 0) \bmod m - 1 & \text{if } k = 0 \\ m - 1 & \text{if } k = 2 \\ m - 2 & \text{if } k = 1 \wedge \\ & SSN(\gamma, 0) = SSN(\gamma, 1) \bmod m - 1 \\ SSN(\gamma, 1) \bmod m - 2 & \text{otherwise} \end{cases} \quad (1.156)$$

1.2.3 Winner selection and claim validation

Definition 28 – Prefixed hash .

Define the hash prefixing function $prefix(H, p)$ as a function which when applied to a hash function H and a constant byte slice p outputs a hash function which for every input returns the hash of the input prefixed by p using H .

$$prefix : (byte* \rightarrow Segment) \times byte* \rightarrow (byte* \rightarrow Segment) \quad (1.157)$$

$$prefix(H, p) : byte* \rightarrow Segment \quad (1.158)$$

$$prefix(H, p)(b) \stackrel{\text{def}}{=} H(p \oplus b) \quad (1.159)$$

Exceptionally, we define the prefixed version of BMT hash (denoted as $BMT[]$) as one that uses the prefixed version of its base hash:

$$BMT[] : byte* \rightarrow Chunk \times byte\{8\} \rightarrow Segment \quad (1.160)$$

$$BMT[p] \stackrel{\text{def}}{=} BMT[prefix(H, p)] \quad (1.161)$$

Definition 29 – Transformed chunk reserve sample.

Let us now define the chunk transformation function $\Delta^C(p)$ for a random nonce prefix p as follows:

$$\Delta^C : \text{Nonce} \rightarrow \text{Chunks} \rightarrow \text{Segment} \quad (1.162)$$

$$\Delta^C(p) : \text{Chunks} \rightarrow \text{Segment} \quad (1.163)$$

$$\Delta^C(p)(c) \stackrel{\text{def}}{=} \text{BMT}[p](\text{DATA}(c), \text{METADATA}(c)) \quad (1.164)$$

Define the transformed reserve sample $RS^C(\gamma, n)$ for game γ and node n as the first 2^d chunks of the node's reserve at block $\text{BLOCK}(\gamma)$, using the ordering defined (see lemma 20 and definition 22) by the hash (see definition 28) of their data using BMT with 256-bit Keccak prefixed with random nonce p as its base hash.

$$RS^C : \Gamma \times \text{Nodes} \rightarrow \text{Chunks}^* \quad (1.165)$$

$$RS^C(\gamma, n) \stackrel{\text{def}}{=} \text{Sampler}(\Delta^C(p), \text{Reserve}(\gamma, n), 2^d) \quad (1.166)$$

where

$$d = \text{SAMPLE_DEPTH} \text{ (see ??)} \quad (1.167)$$

$$p = \text{RSS}(\gamma) \text{ (see definition 27)} \quad (1.168)$$

Definition 30 – Chunk reserve sample commitment hash .

Define $CH^C(\gamma, n,)$ for game γ and node n as the BMT chunk hash of the packed address chunk (see definition 11) packing the chunks of the transformed reserve sample (see definition 29) for game γ and node n .

$$CH^C : \Gamma \times \text{Nodes} \rightarrow \text{Address} \quad (1.169)$$

$$CH^C(\gamma, n) \stackrel{\text{def}}{=} \text{Address}(\text{PAC}^C(RS^C(\gamma, n), p)) \quad (1.170)$$

where

$$p = \text{RSS}(\gamma) \text{ (see definition 27)} \quad (1.171)$$

and where

$$\text{PAC}^C : \text{Chunks}^* \times \text{Nonce} \rightarrow \text{Chunks} \quad (1.172)$$

$$\text{PAC}^C(C, p) \stackrel{\text{def}}{=} \text{CAC} \left(\bigoplus_{i=0}^{\text{len}(C)-1} \text{Seg}(C[i], p), m \right) \quad (1.173)$$

where

$$\text{Seg}(C[i], p) = \text{Address}(C[i]) \oplus \Delta^C(p)(C[i]) \quad (1.174)$$

$$m = LE^{64}(2 \cdot 32 \cdot \text{len}(C)) \quad (1.175)$$

Definition 31 – Transformed slots reserve sample.

Let us now define the slots transformation function $\Delta^S(p)$ for a random nonce p as follows:

$$\Delta^S : \text{Nonce} \rightarrow \text{Chunks} \rightarrow \text{Segment} \quad (1.176)$$

$$\Delta^S(p) : \text{Chunks} \rightarrow \text{Segment} \quad (1.177)$$

$$\Delta^S(p)(c) \stackrel{\text{def}}{=} H[p](\text{Slot}(\text{Stamp}(c))) \quad (1.178)$$

Define the transformed reserve sample $RS^S(\gamma, n)$ for game γ and node n as the first 2^d chunks of the node's reserve at block $\text{BLOCK}(\gamma)$, using the ordering defined by the slot transformation function using prefix p :

$$RS^S : \Gamma \times \text{Nodes} \rightarrow \text{Chunks}^* \quad (1.179)$$

$$RS^S(\gamma, n) \stackrel{\text{def}}{=} \text{Sampler}(\Delta^S(p), \text{Reserve}(\gamma, n), 2^d) \quad (1.180)$$

where

$$d = \text{SAMPLE_DEPTH (see ??)} \quad (1.181)$$

$$p = \text{RSS}(\gamma) \text{ (see definition 27)} \quad (1.182)$$

Definition 32 – Transformed slots reserve sample commitment hash .

Define $CH^S(\gamma, n)$ for game γ and node n as the BMT chunk hash of the packed address chunk (see definition 11) the chunks of the transformed reserve sample (see definition 29) for game γ and node n .

$$CH^S : \Gamma \times \text{Nodes} \rightarrow \text{Address} \quad (1.183)$$

$$CH^S(\gamma, n) \stackrel{\text{def}}{=} \text{Address}(\text{PAC}(RS^S(\gamma, n), m)) \quad (1.184)$$

where

$$m = LE^{64}(32 \cdot 2^d) \quad (1.185)$$

$$d = \text{SAMPLE_DEPTH (see ??)} \quad (1.186)$$

Definition 33 – Weighted selection.

We define $\text{WeightedSelect}(w, k)$ as a sampler function which selects an index $0 < i < \text{len}(w)$ determined by the input nonce (pseudorandom number) k in such a way that the indexes have a probability of being selected proportional to the weights in w .

$$\text{WeightedSelect} : \text{uint256}^+ \times \text{uint256} \rightarrow \text{uint256}^+ \quad (1.187)$$

$$\text{WeightedSelect}(w, k) \stackrel{\text{def}}{=} \begin{cases} i & \text{if } k < w[i] \bmod W[i] \\ \text{WeightedSelect}(w[:i], k) & \text{otherwise} \end{cases} \quad (1.188)$$

such that $i = \text{len}(w) - 1$, and where W (cumulative weights) is defined as

$$W : \text{uint256} + \rightarrow \text{uint256} \quad (1.189)$$

$$W[i] \stackrel{\text{def}}{=} \begin{cases} w[0] & \text{if } i = 0 \\ W[i-1] + w[i] & \text{otherwise} \end{cases} \quad (1.190)$$

Definition 34 – Truth selection.

We determine the truth from reveals through selection weighted by stake density using the truth selection nonce as random input.

$$\text{Truth} : \Gamma \rightarrow \text{RevealEntry} \quad (1.191)$$

$$\text{Truth}(\gamma) \stackrel{\text{def}}{=} R(\gamma)[\text{WeightedSelect}(\text{weights}, \text{TSN}(\gamma))] \quad (1.192)$$

where weights are stake densities such that

$$\text{weights}[i] = \text{Stake}(\gamma, \text{NODE}(R[i])) \cdot 2^{\text{DEPTH}(R[i])} \quad (1.193)$$

where R is the reveals of the round sorted by index:

$$R = \overrightarrow{\text{Seq}}(\text{INDEX}, \text{Reveals}(\gamma)) \quad (1.194)$$

Definition 35 – Honest reveals.

We define honest reveals as the subset of reveals for the round agreeing with the truth in reserve commitment hashes and storage depth.

$$\text{HonestReveals} : \Gamma \rightarrow \text{RevealEntry}^* \quad (1.195)$$

$$\text{HonestReveals}(\gamma) \stackrel{\text{def}}{=} \overrightarrow{\text{Seq}}(\text{INDEX}, \{r \in \text{Reveals}(\gamma) \mid \text{Honest}(r)\}) \quad (1.196)$$

where

$$\text{Honest} : \text{RevealEntry} \rightarrow \{\text{T}, \text{F}\} \quad (1.197)$$

$$\text{Honest}(r) \leftrightarrow \quad (1.198)$$

$$\text{CHC}(r) = \text{CHC}(\text{truth}(\gamma)) \wedge \quad (1.199)$$

$$\text{CHS}(r) = \text{CHS}(\text{truth}(\gamma)) \wedge \quad (1.200)$$

$$\text{DEPTH}(r) = \text{DEPTH}(\text{truth}(\gamma)) \quad (1.201)$$

Definition 36 – Winner selection.

We determine the winner from honest reveals through selection weighted by stake using the winner selection nonce as random input.

$$\text{Winner} : \Gamma \rightarrow \text{RevealEntry} \quad (1.202)$$

$$\text{Winner}(\gamma) \stackrel{\text{def}}{=} \text{HonestReveals}(\gamma)[\text{WeightedSelect}(\text{weights}, \text{WSN}(\gamma))] \quad (1.203)$$

where weights are stakes such that

$$weights[i] = Stake(\gamma, \text{NODE}(HonestReveals[i])) \quad (1.204)$$

1.2.4 Proofs of reserve

Definition 37 – Proof of reserve .

Proof of reserve provides evidence that the reserve is replicating relevant content and shows a proof of recency of retaining chunk data in full integrity.

$$POR : SIP^2 \times Stamps \quad (1.205)$$

$$\Pi^R : \Gamma \times Chunks \times Chunks+ \times \{0, 1, 2\} \rightarrow POR \quad (1.206)$$

$$\Pi^R(\gamma, c, C, k) \stackrel{\text{def}}{=} \langle \quad (1.207)$$

$$\text{WITNESS PROOF} \quad \Pi^{\text{SIP}}(c, d \cdot i), \quad (1.208)$$

$$\text{RETENTION PROOF} \quad \Pi^{\text{SIP}}(C[i], j), \quad (1.209)$$

$$\text{POSTAGE STAMP} \quad Stamp(C[i]) \quad (1.210)$$

$$\rangle \quad (1.211)$$

where

$$i = \mathcal{W}(\gamma, \text{len}(C), k) \quad (1.212)$$

$$j = SSN(\gamma, k) \bmod 128 \quad (1.213)$$

$$d = \frac{SegCnt(c)}{\text{len}(C)} \left(= \begin{cases} 1 & \text{if } C = RS^S \\ 2 & \text{if } C = RS^C \end{cases} \right) \quad (1.214)$$

Definition 38 – Proof of reserve validation.

$$\mathbb{V}^R : \Gamma \times Nodes \times Address \times \{0, 1, 2\} \times POR \rightarrow \{\mathbf{T}, \mathbf{F}\} \quad (1.215)$$

$$\mathbb{V}^R(\gamma, n, ch, k, \pi) \leftrightarrow \quad (1.216)$$

$$\text{RELEVANCE} \quad \mathbb{V}^{\text{SIP}}(ch, d \cdot i, p_w) \wedge \quad (1.217)$$

$$\mathbb{V}^{\text{STAMP}}(ps, \gamma, n) \wedge \quad (1.218)$$

$$\text{RETENTION} \quad \text{DATA}(p_w) = a \wedge \quad (1.219)$$

$$\mathbb{V}^{\text{SIP}}(a, j, p_r) \wedge \quad (1.220)$$

$$\text{RECENCY} \quad i = \mathcal{W}(\gamma, SegCnt(p_w), k) \wedge \quad (1.221)$$

$$j = SSN(\gamma, k) \bmod 128 \wedge \quad (1.222)$$

$$\text{RETRIEVABILITY} \quad PO(a, n) \geq sd \quad (1.223)$$

where

$$a = \text{ADDRESS}(ps) \quad (1.224)$$

$$\pi = \langle p_w, p_r, ps \rangle \quad (1.225)$$

$$sd = \text{DEPTH}(\text{Reveal}(\gamma, n)) \quad (1.226)$$

$$d = \frac{\text{SegCnt}(p_w)}{2^D} \left(= \begin{cases} 1 & \text{if } C = RS^S \\ 2 & \text{if } C = RS^C \end{cases} \right) \quad (1.227)$$

$$D = \text{SAMPLE_DEPTH} \text{ (see ??)} \quad (1.228)$$

Definition 39 – Proof of chunk density validation.

Define $\mathbb{V}^{\text{CD}}(\gamma, p_0, p_1, p_2)$ as the validation function for the proof of chunk density for round γ and proof of reserve and segment inclusion proof pairs p_0, p_1, p_2 .

$$PORT : POR \times SIP \quad (1.229)$$

$$\mathbb{V}^{\text{CD}} : \Gamma \times PORT^3 \rightarrow \{\text{T}, \text{F}\} \quad (1.230)$$

$$\mathbb{V}^{\text{CD}}(\gamma, \pi_0, \pi_1, \pi_2) \leftrightarrow \quad (1.231)$$

$$\text{RIGHT DATA SEGMENT} \quad \text{DATA}(pr_k) = \text{DATA}(pt_k) \text{ for } k \in \{0, 1, 2\} \wedge \quad (1.232)$$

$$\text{RIGHT ADDRESS} \quad \text{SISTER}(pw_k, 0) = ta_k \text{ for } k \in \{0, 1, 2\} \wedge \quad (1.233)$$

$$\text{RIGHT ORDER} \quad ta_0 < ta_1 < ta_2 \wedge \quad (1.234)$$

$$\text{RIGHT SIZE} \quad ta_2 \leq \text{MAX_SAMPLE_VALUE} \quad (1.235)$$

where for $k \in \{0, 1, 2\}$

$$ta_k = H_{\Pi}(p, j_k, pt_k) \quad (1.236)$$

$$j_k = \text{SSN}(\gamma, k) \bmod 128 \quad (1.237)$$

$$\pi_k = \langle \langle pw_k, pr_k, - \rangle, pt_k \rangle \quad (1.238)$$

and where

$$p = \text{RSS}(\text{Prev}(\gamma)) \text{ (see definition 27)} \quad (1.239)$$

Definition 40 – Proof of stamp density validation.

Define $\mathbb{V}^{\text{SD}}(\gamma, ps_0, ps_1, ps_2)$ as the validation function for the proof of stamp density for round γ and postage stamps ps_0, ps_1, ps_2 .

$$\mathbb{V}^{\text{SD}} : \Gamma \times \text{Stamps}^3 \rightarrow \{\text{T}, \text{F}\} \quad (1.240)$$

$$\mathbb{V}^{\text{SD}}(\gamma, ps_0, ps_1, ps_2) \leftrightarrow \quad (1.241)$$

$$\text{RIGHT ORDER} \quad ta_0 < ta_1 < ta_2 \wedge \quad (1.242)$$

$$\text{RIGHT SIZE} \quad ta_2 \leq \text{MAX_SAMPLE_VALUE} \quad (1.243)$$

where for $k \in \{0, 1, 2\}$

$$ta_k = H(Slot(ps_k) \oplus p) \quad (1.244)$$

and where

$$p = RSS(Prev(\gamma)) \text{ (see definition 27)} \quad (1.245)$$

Definition 41 – Proof of entitlement.

Proof of entitlement captures all the evidence a node needs to submit with their claim transaction to validate.

$$POE : PORT^3 \times POR^3 \quad (1.246)$$

$$\Pi^{\text{ENT}} : \Gamma \times Nodes \rightarrow POE \quad (1.247)$$

$$\Pi^{\text{ENT}}(\gamma, n) \stackrel{\text{def}}{=} \langle \quad (1.248)$$

$$\langle \Pi^{\text{POR}}(\gamma, n, PAC^C(crs, p), crs, 0), pt_0 \rangle, \quad (1.249)$$

$$\langle \Pi^{\text{POR}}(\gamma, n, PAC^C(crs, p), crs, 1), pt_1 \rangle, \quad (1.250)$$

$$\langle \Pi^{\text{POR}}(\gamma, n, PAC^C(crs, p), crs, 2), pt_2 \rangle, \quad (1.251)$$

$$\Pi^{\text{POR}}(\gamma, n, PAC(srs), srs, 0), \quad (1.252)$$

$$\Pi^{\text{POR}}(\gamma, n, PAC(srs), srs, 1), \quad (1.253)$$

$$\Pi^{\text{POR}}(\gamma, n, PAC(srs), srs, 2), \quad (1.254)$$

$$\rangle \quad (1.255)$$

where for $k \in \{0, 1, 2\}$

$$pt_k = \Pi_{prefix}^{\text{SIP}}(p, r[i_k], j_k) \quad (1.256)$$

$$j_k = SSN(\gamma, k) \bmod 128 \quad (1.257)$$

$$i_k = \mathcal{W}(\gamma, length(crs), k) \quad (1.258)$$

and where

$$crs = RS^C(\gamma, n) \quad (1.259)$$

$$srs = RS^S(\gamma, n) \quad (1.260)$$

$$p = RSS(Prev(\gamma)) \text{ (see definition 27)} \quad (1.261)$$

Definition 42 – Winner's claim validation.

Define $\mathbb{V}^{\text{POE}}(\gamma, p)$ as the validation function for the proof of entitlement p as part of the

winning claim for game γ :

$$\mathbb{V}^{\text{POE}} : \Gamma \times \text{POE} \rightarrow \{\text{T}, \text{F}\} \quad (1.262)$$

$$\mathbb{V}^{\text{POE}}(\gamma, p) \Leftrightarrow \quad (1.263)$$

$$\text{RESERVE:} \quad (1.264)$$

$$\text{CHUNKS} \quad \forall k \in \{0, 1, 2\}, \mathbb{V}^{\text{R}}(\gamma, n, \text{CHC}(r), k, \pi_k) \wedge \quad (1.265)$$

$$\text{STAMPS} \quad \forall k \in \{0, 1, 2\}, \mathbb{V}^{\text{R}}(\gamma, n, \text{CHS}(r), k, \phi_k) \wedge \quad (1.266)$$

$$\text{RESERVE SIZE:} \quad (1.267)$$

$$\text{CHUNK DENSITY} \quad \mathbb{V}^{\text{CD}}(\gamma, \langle \pi_0, pt_0 \rangle, \langle \pi_1, pt_1 \rangle, \langle \pi_2, pt_2 \rangle) \wedge \quad (1.268)$$

$$\text{STAMP DENSITY} \quad \mathbb{V}^{\text{SD}}(\gamma, \text{PS}(\phi_0), \text{PS}(\phi_1), \text{PS}(\phi_2)) \quad (1.269)$$

where

$$n = \text{NODE}(r) \quad (1.270)$$

$$r = \text{Winner}(\gamma) \quad (1.271)$$

$$p = \langle \langle \pi_0, pt_0 \rangle, \langle \pi_1, pt_1 \rangle, \langle \pi_2, pt_2 \rangle, \phi_0, \phi_1, \phi_2 \rangle \quad (1.272)$$

Corollary 43 – Outpayment scheme is fair.

Rewarding the pot to randomly selected neighbourhood implements a redistribution scheme that is fair across neighbourhoods.

Proof. With the consensus mechanism we can show that the Nash-optimal strategy of nodes is to follow the protocol and consent on the reserve. On the other hand, the optimal strategy for uploaders is to uniformly distribute chunks across the name space. As a consequence, nodes are expected to have identical storage depth and its variance is independent of being chosen. Long term then relative cumulative outpayments by the redistribution game converge to the fair share.

Secondly, we argue that the mode of selecting the winner is fair within neighbourhoods.

Proof.

Chapter 2

Data types and algorithms

2.1 Built-in primitives

2.1.1 Crypto

This section describes the crypto primitives used throughout the specification. They are exposed as buzz built-in functions. The modules are hashing, random number generation, key derivation, symmetric and asymmetric encryption (ECIES), mining (i.e., finding a nonce), elliptic curve key generation, digital signature (ECDSA), Diffie–Hellman shared secret (ECDH) and Reed-Solomon (RS) erasure coding.

Some of the built-in crypto primitives (notably, sha3 hash, and ECDSA ecrecover) are replicating crypto functionality of the Ethereum VM. These are defined here with the help of ethereum api calls to a smart contract. This smart contract just implements the primitives of "buzz" and only has read methods.

Hashing

The base hash function implements Keccak256 hash as used in Ethereum.

Definition 44 – Hashing.

```
// /crypto 1
2
define function hash @input []byte 3
    ?and/with @suff 4
    return segment 5
as 6
```

```

    ethereum/call "sha3" with @input append= @suff
    on context contracts "buzz"

```

7
8

Random number generation

Definition 45 – Random number generation.

```

    // /crypto

```

1

```

define function random type
    return [@type size]byte

```

2
3
4

Script key derivation

The crypto key derivation function implements `scrypt`

Definition 46 – Script key derivation.

```

    // /crypto

```

1

```

define type salt as [segment size]byte

```

2
3
4

```

define type key as [segment size]byte

```

5
6

```

// params for scrypt key derivation function
// scrypt.key(password, salt, n, r, p, 32) to generate key

```

7
8
9

```

define type kdf
    n int // 262144
    r int // 8
    p int // 1

```

10
11
12
13
14

```

define function scrypt from @password
    with salt
    using kdf
    return key

```

15
16
17
18

Mining helper

This module provides a very simple helper function that finds a nonce that when given as the single argument to a mining function returns true.

Definition 47 – Mining a nonce.

```
// /crypto 1
2
define type nonce as [segment size]byte 3
4
define function mine @f function of nonce return bool 5
as 6
    @nonce = random key 7
    return @nonce if call @f @nonce 8
    self @f 9
```

Symmetric encryption

Symmetric encryption uses a modified blockcipher with 32 byte blocksize in counter mode. The segment keys are generated by hashing the chunk-specific encryption key with the counter and hash that again. This second step is required so that a segment can be selectively disclosed in a 3rd party provable way yet without compromising the security of the rest of the chunk.

The module provides input length preserving blockcipher encryption.

Definition 48 – Blockcipher.

```
// /crypto 1
2
// two-way (en/de)crypt function for segment 3
define function crypt.segment segment 4
    with key 5
    at @i uint8 6
as 7
    hash @key and @i // counter mode 8
    hash // extra hashing 9
    to @segment length // chop if needed 10
    xor @segment // xor key with segment 11
12
// two-way (en/de)crypt function for arbitrary length 13
define function crypt @input []byte 14
```

```

    with key
    return [@input length]byte
as
    @segments = @input each segment size // iterate segments
    of input
    go crypt.segment at @i++ with @key // concurrent crypt
    on segments
    return wait for @segments // wait for results
    join // join (en/de)
    crypted segments

```

Elliptic curve keys

Public key cryptography is the same as in Ethereum, it uses the secp256k1 elliptic curve.

Definition 49 – Elliptic curve key generation.

```

// /crypto
define type pubkey as [64]byte
define type keypair
    privkey [32]byte
    pubkey

define type address as [20]byte

define function address pubkey
    return address
as
    hash pubkey
    from 12

define function generate
    ?using entropy
as
    @entropy = random segment if no @entropy
    http/get "signer/generate?entropy=" append @entropy
    as keypair

```


Asymmetric encryption

Asymmetric encryption implements ECIES based on the secp256k1 elliptic curve.

Definition 50 – Asymmetric encryption.

```
// /crypto 1
2
3
4
5
6
7
8
9
define function encrypt @input []byte
    for pubkey
    return [@input length]byte
define function decrypt @input []byte
    with keypair
    return [@input length]byte
```

Signature

Crypto's built-in signature module implements secp256k1 elliptic curve based ECDSA. The actual signing happens in the external signer running as a separate process (possibly within the secure enclave). As customary in Ethereum, the signature is represented and serialised using the r/s/v format,

Definition 51 – Signature.

```
// /crypto 1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
define type signature
    r segment
    s segment
    v uint8
    signer private keypair
define type doc
    preamble []byte
    context []byte
    asset segment
define function sign @input []byte
    by keypair
    return signature
as
```

```

@doc = doc{ "swarm signature", context caller, @input }      19
@sig = http/get "signer/sign?text=" append @doc              20
      append "&account=" append @keypair pubkey address      21
      as signature                                           22
@sig signer = @keypair                                       23
@sig                                                         24
                                                         25
define function recover signature                             26
  with @input []byte                                         27
  from @caller []byte                                         28
  return pubkey                                               29
as                                                         30
  @doc = doc{ "swarm signature", @caller, @input } as bytes  31
  ethereum/call "ecrecover" with                             32
    on context contracts "buzz"                               33
    as pubkey                                                 34

```

Diffie-Hellmann shared secret

The shared secret module implements elliptic curve based Diffie–Hellmann shared secret (ECDH) using the usual secp256k1 elliptic curve. The actual DH comes from the external signer which is then hashed together with a salt.

Definition 52 – Shared secret.

```

// /crypto                                                    1
                                                         2
define function shared.secret between keypair                 3
  and pubkey                                                  4
  using salt                                                  5
  return [segment size]byte                                   6
as                                                         7
  http/get "signer/dh?pubkey=" append @pubkey append "&      8
    account=" @keypair address
    hash with @salt                                           9

```

Erasure coding

Erasure coding interface provides wrappers `extend/repair` for the encoder/decoder that work directly on a list of chunks.¹

Assuming n out of m coding. `extend` takes a list of n data chunks and an argument for the number of required parities. It returns the parity chunks only. `repair` takes a list of m chunks (extended with *all* parities) and an argument for the number of parities $p = m - n$, that designate the last p chunks as parity chunks. It returns the list of n repaired data chunks only. The encoder does not know which parts are invalid, so missing or invalid chunks should be set to `nil` in the argument to repair. If parity chunks are needed to be repaired, you call `repair @chunks with @parities`; `extend` with `@parities`

Definition 53 – CRS erasure code interface definition.

```
// /crypto/crs 1
2
define function extend @chunks [] chunk 3
  with @parities uint 4
  return [@parities] chunk 5
6
define function repair @chunks [] chunk 7
  with @parities uint 8
  return [@chunks length - @parities] chunk 9
```

Definition 54 – CRS erasure coding parameters.

```
// /crypto/crs 1
2
define strategy as "race"|"fallback"|"disabled" 2
3
define type params 4
  parities uint 5
  strategy 6
```

2.1.2 State store

Definition 55 – State store.

```
// /statestore 1
2
define type key [] byte 3
```

¹Cauchy-Reed-Solomon erasure codes based on <https://github.com/klauspost/reedsolomon>.

```

define type db []byte 4
define type value []byte 5

define function create db 6
7
8
define function destroy db 9
10
define function put value 11
    to db 12
    on key 13
14
define function get key 15
    from db 16
    return value 17

```

2.1.3 Local context and configuration

Definition 56 – Context.

```

// /context 1
2
define type contract as "buzz"|"chequebook"|"postage"|" " 3
4
define type context 5
    contracts [contract]ethereum/address 6

```

2.2 Bzz address

2.2.1 Overlay address

Swarm's overlay network uses 32-byte addresses. In order to help uniform utilisation of the address space, these addresses must be derived using a hash function. A Swarm node must be associated a Swarm base account or bzz account, which is an Ethereum account that the node (operator) must possess the private key for. The node's overlay address is derived the public key of this account.

Definition 57 – Swarm overlay address of node A.

$$\text{overlayAddress}(A) \stackrel{\text{def}}{=} \text{Hash}(\text{ethAddress}|\text{bzzNetworkID}) \quad (2.1)$$

where

- *Hash* is the 256-bit Keccak SHA3 hash function
- *ethAddress* - the ethereum address (bytes, not hex) derived from the node's base account public key: $\text{account} \stackrel{\text{def}}{=} \text{PubKey}(K_A^{\text{bzz}})[12 : 32]$, where
 - *PubKey* is the *uncompressed* form of the public key of a keypair *including* its 04 (uncompressed) prefix.
 - K_A^{bzz} refers to the node's bzz account key pair
- *bzzNetworkID* is the bzz network id of the swarm network serialised as a little-endian binary *uint64*.

In a way, deriving the node address from a public key means the overlay address space is not freely available: to occupy an address you must possess the private key of the address which other nodes need to verify. Such authentication is easy using a digital signature of a shared consensual piece of text, see 2.2.

2.2.2 Underlay address

To enable peers to locate the a node on the network, the overlay address is paired with an underlay address. The underlay address is a string representation of the node's network location on the underlying transport layer. It is used by nodes to dial other nodes to establish keep-alive peer to peer connections.

2.2.3 BZZ address

Bzz address is functionally the pairing of overlay and underlay addresses. In order to ensure that an overlay address is derived from an account the node possesses as well as verifiably attest to an underlay address a node can be called on, bzz addresses are communicated in the following transfer format:

Definition 59 – Swarm bzz address transfer format.

```
// ID: /swarm/handshake/1.0.0/bzzaddress 1
syntax = "proto3"; 2
3
```

```

message BzzAddress {
    bytes Underlay = 1;
    Signature Sig = 2;
    bytes Overlay = 3;
}

```

Here the signature is attesting to the association of an overlay and an underlay address for a network.

Definition 60 – Signed underlay address of node A.

$$\text{signedUnderlay}(A) \stackrel{\text{def}}{=} \text{Sign}(\text{underlay}|\text{overlay}|\text{bzzNetworkID}) \quad (2.2)$$

of the underlay with overlay and bzz network ID appended as plaintext and hashes the resulting public key together with the bzz network ID.

Definition 61 – Node addresses: overlay, underlay, bzz address.

```

// /bzz
define type overlay as [segment size]byte
define type underlay []byte

define function overlay.address of pubkey
    within @network uint64
as
    hash @pubkey address and @network
    as overlay

define function valid bzz.address
    within @network uint64
as
    assert @bzz.address overlay == overlay.address of crypto/
        recover from @bzz.address signature with @bzz.address
        @underlay
        within @network

define function bzz.address of overlay
    from underlay

```

```

    by @account ethereum/address 20
as 21
    @sig = crypto/sign @underlay by @account 22
    bzz.address{ @overlay, @underlay, @sig } 23

```

In order to get the overlay address from the transfer format peer info, one recovers from signature the peer's base account public key using the plaintext that is constructed as per 60. From the public key, the overlay can be calculated as in 3. The overlay address thus obtained needs to be checked against the one provided in the handshake.

Signing of the underlay enables preflight authentication of the underlay of a trusted but not connected node.

Since underlays are meant to be volatile, we can assume and in fact expect multiple underlays signed by the same node. However, these are meant to be temporally ordered. So one with a newer timestamp invalidates the older one.

In order to make sure that the node connected through that underlay does indeed operate the overlay address, its authentication must be obtained through the peer connection that was initiated by dialing the underlay. This protects against malicious impersonation of a trusted overlay potentially.

2.3 Chunks, encryption and addressing

2.3.1 Content addressed chunks

First let us define some basic types, such as *payload*, *span*, *segment*. These fixed length byte slices enables verbose expression of fundamental units like `segment size` or `payload size`.

Definition 62 – `segment`, `payload`, `span`, `branches`.

```

// /chunk 1
2
define type segment as [32]byte // unit for type definitions 3
define type payload as [:4096]byte // variable length max 4 4
    Kilobyte
define type span as uint64 // little endian binary 5
    marshalled 6
define function branches 7

```

```
as payload size / segment size
```

8

Now let's turn to the definition of *address*, *key* and *reference*:

Definition 63 – Chunk reference.

```
// /chunk

define type address as [segment size]byte

define type reference
    address // result of bmt hash
    key if context.encryption // decryption key optional (
        context dependent)
```

1

2

3

4

5

6

7

Now, define chunk as a object with span and payload.

Definition 64 – Content addressed chunk.

```
// /chunk

define type chunk
    span // length of data span subsumed under node
    payload // max 4096 bytes

define function address of chunk
as
    @chunk payload bmt/hash with @chunk span

define function create from payload
    ?over span
as
    @span = @payload length if no @span
    @chunk = chunk{ @span, @payload }
    return @chunk if no context encryption
    @key = encryption.key for @chunk
    @chunk encrypt with @key
```

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

17

18

Where length is the content of the length field and reference size is the sum of size of the referencing hash value and that of the decryption key, which is currently 64, as we use 256-bit hashes and 256-bit keys.

In order to remove the padding after decryption before returning the plaintext chunk.

Definition 65 – Span to payload length.

```
// /chunk 1
2
define function payload.length of span 3
as 4
    while @span >= 4096 5
        @span = @span + 4095 6
        / 4096 7
        * reference size 8
    return @span 9
```

Finally, we can define the public API of chunks for retrieval and storage.

Definition 66 – Chunk API retrieval.

```
// /chunk 1
2
define function retrieve reference 3
has api GET on "chunk/<reference>" 4
as 5
    retrieve @reference address as chunk 6
    (decrypt with @reference key if @reference key) 7
```

Definition 67 – Chunk API: storage.

```
// /chunk 1
2
3
define function store payload 4
    ?over span 5
has api POST on "chunk/(?span=<span>)" 6
    from payload as body 7
as 8
    @chunk = create from @payload over @span 9
    reference{ @chunk address, @chunk key } 10
```

2.3.2 Single owner chunk

Single owner chunks are the second type of chunk in swarm. They constitute the basis of swarm feeds.

Definition 68 – Single owner chunks.

```
// /soc 1
2
// data structure for single owner chunk 3
define type soc 4
  id [segment size]byte // id 'within' owner 5
  namespace
  signature crypto/signature // owner attests to <content, 6
  id>
  chunk // content: embeds a content 7
  chunk
8
// constructor for single owner chunks 9
define function create from chunk 10
  by @owner crypto/keypair 11
  on @id [segment size]byte 12
as 13
  @sig = crypto/sign @id and @chunk address by @owner 14
  soc{ @id, @sig, @chunk } 15
16
define function address soc 17
as 18
  hash @soc id and @soc signature signer address 19
```

Definition 69 – Single owner chunk API: retrieval.

```
// /soc 1
2
define function retrieve @id [segment size]byte 3
  by @owner ethereum/address 4
  ?with key 5
has api GET on "soc/<owner>/<id>(?key=<key>)" 6
as 7
  retrieve hash @id and @owner 8
  as soc 9
    chunk (decrypt with @key if @key) 10
```

Definition 70 – Single owner chunk API: storage.

```
// /soc 1
2
define function store payload 3
  on @id [segment size]byte 4
  by @owner ethereum/address 5
  ?over span 6
has api POST on "soc/<owner>/<id>?span=<span>&encrypt=<encrypt> 7
"
  from <payload> as body 8
as 9
  @span = @payload length if no @span 10
  @chunk = chunk{ @span, @payload } 11
  if context encryption then 12
    @key = encryption.key for @chunk 13
    @chunk encrypt= with @key 14
  @soc = create from @chunk on @id by private key of @owner 15
  reference{ @soc store address, @key } 16
```

2.3.3 Binary Merkle Tree Hash

The hashing method used to obtain the address of the default content addressed chunk is called the binary Merkle tree hash, or BMT hash for short.

Calculating the BMT hash

The base segments of the binary tree are subsequences of the chunk content data. The size of segments is 32 bytes, which is the digest size of the *base hash* used to construct the tree. Given the Swarm hash tree used to represent files (see 2.5.1) assumes that intermediate chunks package references to other chunks.

Obtaining the BMT hash of a sequence involves the following steps:

1. *padding* - If the content is shorter than the maximum chunk Size (4096 bytes), it is padded with zeros up to chunk size. Note that this zero padding is only for hashing and does not impact chunk data sizes.
2. *chunk data layer* - Calculate the base hash of *pairs of segments* in the padded chunk, i.e., segment size (2 * 32) units of data and concatenate the results.

3. *building the tree* - Repeat previous step on the result until the result is just one section.
4. *calculate span* - Calculate the span of the data, i.e., the size of the data that is subsumed under the chunk represented by the unpadded data as a 64-bit little-endian integer value (see 2.5.1).
5. *integrity protection* - Prepend the span to the root hash of the binary tree and calculate the *base hash* of the data.

Definition 71 – BMT hash.

```
// /bmt 1
2
define function hash payload 3
  with span 4
as 5
  @padded = @payload as [:chunk size]byte // use zero 6
             padding
  // for BMT hashing only 7
  hash @span and root of @padded over chunk size 8
9
define function root of @section []byte 10
  over @len uint 11
as 12
  return hash @section // data level 13
    if @len == 2 * segment size 14
  @len /= 2 // recursive call 15
  @children = @section each @len go self over @len 16
  wait for @children 17
  join hash 18
```

Inclusion proofs

Having the segments align with the hashes packaged in these chunks one can extend the notion of inclusion proofs to files. The BMT hash enables compact 3rd party verifiable segment inclusion proofs.

2.3.4 Encryption

Symmetric encryption in Swarm is using a slightly modified version blockcipher in counter mode.

The encryption seed for the chunk is derived from the master seed if given, otherwise just

generated randomly.

The reference to a single chunk (and the whole content) is the concatenation of the hash of encrypted data and the decryption key (see 63). This means the encrypted Swarm reference (64 bytes) will be longer than the unencrypted one (32 bytes). When a node syncs encrypted chunks, it does not share the full references (or the decryption keys) with the other nodes in any way. Therefore, other nodes will be unable to access the original data, or in fact, even to detect whether a chunk is encrypted.

Definition 72 – Chunk encryption/decryption API.

```
// generate key for a chunk
define encryption.key for chunk
  ?with @seed [segment.size] byte
as
  return crypto/random key if no @seed // generate new
  hash @seed and @chunk address

define function encrypt chunk
  with key
as
  @segments = @chunk data pad to chunk size
  each segment size
    go crypt at @i++ with @key
  @span = chunk span crypt at branches with @seed
  @payload = wait for @segments
  join
  chunk{ @span, @payload }

define function decrypt chunk
  with key
as
  @span = @chunk span
  crypt at branches with @key
  @segments = chunk data to @span payload.length
  each segment size
    go crypt at @i++ with @key
  @payload = wait for @segments
  join
  chunk{ @span, @payload }
```

Encrypted Swarm chunks are not different from plaintext chunks and therefore no change is needed on the P2P protocol level to accommodate them. The proposed encryption scheme is end-to-end, meaning that encryption and decryption is done on endpoints, i.e., where the http proxy layer runs. This has an important consequence that public gateways cannot be used for encrypted content. On the other hand, the apiary modular design allows for client side encryption on top of external APIs while proxying all other calls via the gateway.

2.4 Postage stamps

2.4.1 Witness type

There can be different implementations of postage stamps that differ in the structure and semantics of the *proof of payment*. To allow for new cryptographic mechanisms to be used as they are developed, the `witnessType` argument indicates the type of the witness used.

Witness type 0 stands for ECDSA witness, which is an ECDSA signature on the byte slice resulting from the concatenation of 1) preamble constant 2) chunk hash 3) batch reference 4) valid until date.² This is the bare minimum that postage stamp contracts and clients must implement.³

Witness type 1 refers to the RSA witness, which is an RSA signature on the same 128 bytes as above. The binary encoding of the RSA signature is of variable length, and is an Solidity ABI encoded array of the RSA signature s .⁴

The RSA witness is specified so that blind stamping services can be implemented in a simple fashion, in order to mitigate the privacy issues arising from the ability to link chunks signed with the same private key. Even though blind ECDSA signatures also exist, their protocol requires more rounds of communication, making the implementation of such a service more complex, more error-prone and less performant.

The inclusion of the entire public key in each RSA witness rather than storing the public key in contract state and just referencing it from the witness is justified by reducing the gas costs

²The binary encoding of the ECDSA signature is 65 bytes resulting from the concatenation of the r (32 bytes), s (32 bytes) and v (1 byte) parameters of the signature, in this order. The signature is calculated on the secp256k1 elliptic curve, just like the signatures of Ethereum transactions.

³The ECDSA witness is the simplest and cheapest solution both in terms of gas consumed by the stamp verification contract and in terms of computational resources used off chain. Also, it does not rely on cryptographic assumptions in addition to those on which Ethereum critically relies, therefore as long as Ethereum is considered cryptographically secure, no advance in cryptography can render this witness type insecure. This is the justification for this witness type to be the only mandatory witness type to be implemented.

⁴as defined in *PKCS #1*, <https://tools.ietf.org/html/rfc8017> and the RSA public key parameters n (RSA modulus) and e (public exponent).

of interactions with the contract as well as future-proofing the design in case contract state rent is introduced in Ethereum. These considerations are more important than the brevity of postage stamps, marginally reducing the bandwidth costs of uploading and forwarding stamped content.

Note that cryptographic advances can render RSA witnesses insecure without rendering Ethereum insecure, therefore RSA witnesses can be phased out in future versions of the protocol, if the security of RSA signatures gets compromised. Note, furthermore, that such blind signing services are not entirely trustless, through the damage they can incur is bounded. Trustless blind stamping services based on ZK proofs are not feasible at this stage, as the current algorithms are not sufficiently performant for the purpose, but given the rapid advances in the field, the development of suitable algorithms can be expected in the future, in which case a corresponding witness type will have to be specified in a separate SWIP.

2.4.2 Contract Upgrades

In order to facilitate the upgrade of the contract either in case of a discovered vulnerability or some feature extension (such as adding new witness types), it is recommended that the part holding the funds with the database of payments and the part that verifies witnesses are in separate contracts so that a backwards-compatible upgrade can be performed with minimal disruption.

In order to avoid centralized control, it is also recommended that it is the witness-verifying contract that is referenced in client configuration so that client operators can independently decide for themselves when and whether to switch to a new contract, as they become available.

Nodes participating in the same postage system are configured to reference the same contract on the same blockchain. This contract must conform to the following interface:

Definition 73 – Postage contract.

This accessor method returns `true` if the proof embodied by `witness` checks out for all other arguments within the claimed validity period, i.e. when `block.timestamp` (the output of `TIMESTAMP` EVM opcode) is between `beginValidity` (inclusive) and `endValidity` (exclusive). Outside of the validity period, the return value is `undefined`.

Definition 74 – Postage stamp basic types: `batchID`, `address`, `witness`, `stamp`, `validity`.

```
// /postage 1
2
define type batchid as [segment size]byte 3
```

```

define type address as bzz/address
define type witness as crypto/signature

// postage stamp
define type stamp
    batchid
    address
    witness

define function valid stamp
as
    // check validity on blockchain
    ethereum call "valid" using context contracts "postage"
        with @stamp

```

2.5 Files, manifests and data structures

2.5.1 Files and the Swarm hash

This table gives an overview of data sizes a chunk span represents, depending on the level of recursion.

	span					
	unencrypted			encrypted		
level	chunks	\log_2 of bytes	standard	chunks	\log_2 of bytes	standard
0	1	12	4KB	1	12	4KB
1	128	19	512KB	64	18	256KB
2	16,384	26	67MB	4,096	24	16MB
3	2,097,152	33	8.5GB	262,144	30	1.07GB
4	268.44M	40	1.1TB	16.78M	36	68.7GB
5	34,359.74M	47	140TB	1,073.74M	42	4.4TB

Table 2.1: Size of chunk spans

Calculating the Swarm Hash

Client-side custom redundancy is achieved by RS erasure coding; using it necessitates some RS parameters.

Definition 75 – File API: upload/storage; swarm hash.

```

// /file
1
2
define function encode @levels []chunk stream
3
4   for @level uint
5 as
6   @chunks = @levels at @level // read chunk stream
7   @crs = context crs //
8   @m = branches (- @crs parities if @crs)
9
10  @parent = read @m from @chunks // read up to m chunks from
11  stream blocking
12  (append crs/extend with @crs parities if @crs)
13  each chunk/store // package children
14  reference
15  join as chunk
16
17  if @levels length == @level+1 then
18    @levels append= stream{}
19    go self @levels for @level+1
20
21  write @parent to @levels at @level+1
22  if no @chunks then
23    close @levels at @level + 1
24  else
25    self @chunks for @level
26
27 define function split @data byte stream
28 as
29   @level = chunk stream{}
30   go @data each chunk size as chunk
31   write to @chunks
32   @level
33
34 define function upload byte stream as @data
35 has api POST on "file/" from @data as body
36 as
37   @levels append= @data split //
38   go encode @levels for 0
39   @top = @levels each wait for // wait for all levels to
40   close
41   return @top at 0 // return root hash as

```

address

Definition 76 – File API: download/retrieval.

```
// /file 1

define function copy to @reader stream of byte{} 2
  from @chunks stream of [] chunk 3
  using @buffers stream of [@buffer.size] stream of [branches] 4
  chunk 5
as 6
  @chunks = read @buffers if no @chunks 7
  @chunk = read @chunks 8
  if no @chunk 9
    write @chunks to @buffer 10
    self @reader using @buffers 11
  write @chunk data to @reader 12
13
define function download reference 14
  ?using @buffer.size uint64 15
has api GET on "file/<reference>" 16
as 17
  @reader = stream of byte{} 18
  @buffers = stream of [@buffer.size] stream of [branches] 19
  chunk
  chunk/retrieve @reference // root chunk 20
  retrieval
  go decode into @buffer down to 1 // traverse 21
  copy into @reader from @buffers 22
23
define function decode chunk 24
  into @response chunk stream 25
  ?down to @limit uint8 26
as 27
28
  @crs = context crs 29
  @all = @m = branches 30
  if @crs then 31
    @m -= @crs parities 32
    if @crs strategy is not "race" then 33
      @all = @m 34
35
```

```

@chunks = @chunk segments up to @all      36
    each go as reference retrieve          37
wait for @m in @chunks                     38
                                           39
if @crs then                               40
    cancel @chunks                         41
    @chunks = @chunks crs/repair with @crs parities 42
                                           43
if @chunk span < chunk size exp @limit + 1 then 44
    @chunks each into @reader             45
                                           46
                                           47
@chunks each go self down to @limit       48

```

Definition 77 – File info.

```

// /file                                     1
                                           2
define type info                             3
    mode          int64                     4
    size          int64                     5
    modified      time                      6

```

2.5.2 Manifests

Manifests represent a mapping of strings to references. The primary purpose is to implement document collections (websites) on swarm and enable URL-based addressing of content. This section defines the data structures relevant for manifests as well as the algorithms for lookup and update which implement the manifest API (see ??).

A manifest entry can be conceived of as metadata about a file pointed to and retrievable by its reference (see 2.5.1). The metadata is quite diverse, ranging from information needed for access control, file information similar to one given on file systems, information needed for erasure coding (see ?? and 2.1.1), information for browser, i.e., response headers such as content type (MIME info) and most importantly the reference to the file. Using manifests as simple key-value store is exemplified by access control (see 2.6 and ?? for the specification).

Definition 78 – Manifest entry.

```

// /manifest                                 1
                                           2

```

```

// manifest entry encodes attributes
define type entry
    file/info          // FS file/dir info
    access/params      // access control params
    crs/params         // erasure coding - CRS params
    reference          // reference
    headers             // http response headers

define type headers
    content.type [segment size]byte

```

Definition 79 – Manifest data structure.

```

// /manifest

define type node
    entry *entry          // reference to chunk serialised as
    entry
    forks [<256]fork      // sparse array of max 256 fork

// fork encodes a branch
define type fork
    prefix segment        // compaction
    node *node            // reference to chunk serialised as node

```

Definition 80 – Manifest API: path lookup.

```

// /manifest

define function lookup @path []byte
    in *node
has api GET on "/manifest/<@node:reference>/<@path>"
as
    context access = @node entry access/params
    // manifest is a compacted trie
    @fork = @node forks at head @path
    // if @path empty, the paths matched return the entry
    if no @path then
        return @node entry

    if @fork prefix is prefix of @path then // including ==
        return self @path from @fork prefix length

```

```

        in @fork node
fail with "not found"

```

Definition 81 – Manifest API: update.

```

// /manifest
1
2
define function add *entry
3
    to *node
4
    on @path []byte
5
has api PUT on "/manifest/<@node>"
6
7
as
8
    // if called on nil call on zero value
9
    @node = node{} if no @node
10
11
    // if empty path then change entry field of node
12
    if no @path then
13
        @node entry = @entry
14
        return store @node
15
16
    // lookup the fork based on the first byte of path
17
    @fork = @node forks at head @path
18
    // if no fork yet, add the singleton node
19
    if no @fork then
20
        @node forks at head @path =
21
            fork{@path, store node{@entry}}
22
        return store @node
23
24
    @common = prefix of @path and @fork prefix // common cannot
25
        be empty
26
    @rest = @fork prefix from @common length
27
    @newnode = node{}
28
    @newnode forks at head @rest = fork{@rest, @fork node}
29
    @midnode = self @entry to @newnode on @path from @common
    length
30
    @node forks at head @path = fork{ @common, @midnode }
31
    @node store

```

Definition 82 – Manifest API: remove.

```

// /manifest
1

```

```

define function remove @path []byte
    from *node
has api DELETE on "/manifest/<@node>/<@path>"
as
    // if called on nil call on zero value
    return node{} if no @node

    // if empty path then change entry field of node
    if no @path then
        return nil if @node forks length == 0
        @node entry = nil //entry exists
        return store @node

    // lookup the fork based on the first byte of path
    @fork = @node forks at head @path
    // if no fork yet, add the singleton node
    return @node if no @fork

    @common = prefix of @path and @fork prefix // common cannot
        be empty
    return @node if @common and @fork prefix have different
        length // path not found

    @rest = @fork prefix from @common length
    @newnode = self @rest from @fork node
    if no @newnode then // deleted item was terminal
        node, delete fork
        @node forks at head @res = nil
    else if @newnode forks length == 1 then // compact non-
        forking nodes
        @singleton = @newnode forks first
        @newprefix = @common append @singleton prefix
        @node forks at head @path =
            fork{ @newprefix, @singleton node }
    else
        @node fork at head @path node = @newnode

    @node store

```

Definition 83 – Manifest API: merge.

```

// /manifest
1
2
define function merge @new *node
3
  to @old *node
4
has api POST on "/manifest/<@old:reference>/<@new:reference>"
5
as
6
  // if called on nil call on zero value
7
  return @new if no @old
8
  return @old if no @new
9
  @node = node{ @new or @old }
10
  @new forks pos or @old forks pos
11
  each bit @pos go
12
    @fork = merge.fork @new forks at @pos
13
    to @old forks at @pos
14
    @node forks at @fork prefix head = @fork
15
  @node store
16
17
define function merge.fork @new fork
18
  to @old fork
19
as
20
  @common = prefix of @new prefix and @old prefix
21
  @restnew = @new prefix from @common length
22
  @restold = @old prefix from @common length
23
  if no @restnew and no @restold then
24
    return fork{@common, merge.node @new reference to @old
25
      reference}
26
  @node = add @new reference to nil on @restnew
27
  add to @old reference @restold
28
  fork{ @common, @node }

```

2.5.3 Resolver

Definition 84 – Resolver.

```

// /resolver
1
2
define type resolver
3
  api url
4
  address ethereum/address
5
  tlds []string
6
7

```

```

define function resolve @host []byte through @resolver      8
as                                                         9
    @tld = @host split on '.' last      @resolver = @resolvers 10
    any tlds any == @tld                                                         11
    ethereum/call "getContentTypeHash" of @host using @resolver api 12
    at @resolver address
    as address                                                         13

```

2.5.4 Pinning

Definition 85 – Pinning.

```

// /pin                                                         1

define type pin                                                         2
    reference                                                         3
    chunks uint64                                                         4
    5
    6
define function list                                                         7
    return []pin                                                         8
has api GET "/pin/"                                                         9
    10
define function view reference                                                         11
    return uint64                                                         12
has api GET "/pin/<reference>"                                                         13
    14
define function pin reference                                                         15
    return uint64                                                         16
has api PUT "/pin/<reference>"                                                         17
    18
define function pin reference                                                         19
    return uint64                                                         20
has api DELETE "/pin/<reference>"                                                         21

```

2.5.5 Tags

Definition 86 – Tags.

```
// /tag 1
2
3
define type tag 4
    id [segment size]byte 5
    reference // the current root 6
    complete bool // if local upload finished 7
    total uint64 // number of chunks expected 8
    split uint64 // number of chunks split 9
    stored uint64 // number of chunks stored locally 10
    seen uint64 // number of chunks already in db 11
    sent uint64 // number of chunks sent with push-sync 12
    synced uint64 // number of chunks synced 13
14
define function list 15
    return []tag 16
has api GET "/tag/" 17
18
define function view reference 19
    return uint64 20
has api GET "/tag/<reference>" 21
22
define function add reference 23
    return uint64 24
has api POST "/tag/" 25
26
define function remove reference 27
    return uint64 28
has api DELETE "/tag/<reference>" 29
```

2.5.6 Storage

Definition 87 – Public storage API.

```
// /bzz 1
2
3
define function upload @data stream of byte 4
has api POST on "bzz:./<host>/<path>" from @data as body 5
as 6
    @root = resolver/resolve @host
```

```

    @manifest = access/unlock @root
    @entry = file/upload @data
    @reference = chunk/store @entry as []byte
    manifest/add @reference to @manifest on @path

define function download @path []byte from @host []byte
has api POST on "bzz:<host>/<path>" from @data as body
as
    @root = resolver/resolve @host
    @manifest = access/unlock @root
    @entry = manifest/lookup @path in @manifest
    file/download @entry reference

```

2.6 Access Control

Definition 88 – Access control: auth, hint, parameters, root access manifest.

```

// /access

define type auth as "pass"|"pk"
define type hint as [segment size]byte

// access control parameters
define type params
    auth                // serialises uint8
    publisher crypto/pubkey // 65 byte
    salt                // salt for scrypt/dh
    hint                // hint to link identity
    act *node           // reference to act manifest
    root
    kdf                 // params for scrypt
}

// root access manifest
define type root
    params
    reference

```

Definition 89 – Session key and auth with credentials.

```
// /access 1
2
define function session.key.pass from hint 3
    with salt 4
    using kdf 5
as 6
    crypto/scrypt from input/password using @hint 7
    with @salt using @kdf 8
9
define function session.key.pk from hint 10
    with crypto/pubkey 11
    using salt 12
as 13
    crypto/shared.secret between 14
        input/select key by @hint 15
        and @pubkey 16
    hash with @salt 17
18
define function session.key 19
    using params 20
as 21
    if @params auth == "pass" then 22
        return session.key.pass from @params hint 23
        with @params salt using @params kdf 24
25
    session.key.pk from @params hint 26
    with @params publisher using @params salt 27
```

Definition 90 – Access key.

```
// /access 1
2
define function access.key 2
    using params 3
as 4
    @key = session.key using @params 5
    return @key if no @params act 6
    act.lookup @key in @params act 7
8
define function act.lookup key 9
    in @act *node 10
as 11
```

```

manifest/lookup hash @key and 0      12
    in @act                          13
        xor hash @key and 1          14

```

Definition 91 – Access control API: lock/unlock.

```

// /access      1
2
// control      3
define function lock reference      4
    using params      5
has api POST on "/access/<address>"      6
    with @params as body      7
as      8
    @key = hash @reference address and @key      9
    @encrypted = @reference as bytes      10
        crypto/crypt with @key      11
    root{ @params, @encrypted }      12
        store      13
14
15
define function unlock address      16
has api GET on "access/<address>"      17
as      18
    @root = retrieve @address      19
        try as root otherwise return @address      20
    @key = access.key using @root params      21
    @root encrypted crypto/crypt with @key      22
        as *node      23

```

Definition 92 – ACT manipulation API: add/remove.

```

// /access      1
2
define type act as manifest/node      3
4
define function add @keys []crypto/pubkey      5
    to *root      6
has api PUT on "/access/<root>/"      7
    with @keys as body      8
as      9
    // get params from the root access structure      10

```

```

@params = retrieve @root as root params      11
@access.key = access.key using @params        12
@keys each @key                              13
    @session.key = session.key using @params  14
    manifest/add @access.key xor hash @session.key with 1 15
        to @act on hash @session.key with 0    16
                                                17
                                                18
define function remove @keys [] crypto/pubkey 19
    from *root                                20
has api DELETE on "/access/<root>"            21
    with @keys as body                        22
as                                              23
    // get params from the root access structure 24
    @params = retrieve @root as root params    25
    @keys each @key                          26
        @session.key = session.key using @params 27
        manifest/remove hash @session.key with 0 28
            from @params act                    29

```

2.7 PSS

2.7.1 PSS message

2.7.2 Direct pss message with trojan chunk

Pss has two fundamental types, a message and a trojan chunk structure which wraps the encrypted serialised message and contains a nonce that is mined to make the resulting chunk's content address (BMT hash) to match the targets.

Definition 93 – Basic types: topic, targets, recipient, message and trojan.

```

// /pss 1
2
3
define type topic as [segment size] byte // 4
    obfuscated topic matcher
define type targets as [][] byte // overlay prefixes 5
define type recipient as crypto/pubkey 6
7

```

```

// pss message
define type message
    seal    segment
    payload [! :4030] byte    // varlength padded to 4030B

// trojan chunk
// the nonce
define type trojan
    nonce    segment    // the nonce to mine
    key    pubkey    // compressed format
    message [4064] byte    // encrypted msg

```

The message is encoded in a way that allows integrity checking and at the same time obfuscates the topic. The operation to package the payload with a topic is called *sealing*

Definition 94 – Sealing/unsealing the message.

```

// /pss

define function seal @payload [] byte
    with topic
as
    @seal = hash @payload and @topic // obfuscate topic
    xor @topic
    return message{ @seal, @payload }

define function unseal message
    with topic
as
    @seal = hash @message payload and @topic
    if @topic == @seal xor @message seal then // check
        return @payload
    return nil

```

Functions `wrap/unwrap` transform between message and trojan chunk. `wrap` takes an optional recipient public key to asymmetrically encrypt the message. The targets are a list of overlay address prefixes derived from overlay addresses of recipients, with length specified to guarantee that a chunk matching it will end up with the recipient solely as a result if push-syncing.

Definition 95 – Wrapping/unwrapping.

```

// /pss
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
define function wrap message
  for recipient
  to targets
as
  @msg = @message
    (crypto/encrypt for @recipient if @recipient)

  @nonce = crypto/mine @n such that
    @targets any is prefix of
      trojan{@n, @msg} as chunk address
  trojan{@nonce, @msg} as chunk

define function unwrap chunk
  for recipient
as
  @chunk bytes
    (crypto/decrypt for @recipient if @recipient)
    as message

```

When a chunk arrives at the node, `pss/deliver` is called as a hook by the storage component. First the message is unwrapped using the recipient private key and unsealed with all the topics API clients subscribed to. If the unsealing is successful, message integrity as well as topic matching is proven so the payload is written into the stream registered for the topic in question.

Definition 96 – Incoming message handling.

```

// /pss
1
2
3
4
5
6
7
8
9
10
11
12
// mailbox is a handler type, expects payload
// sent sealed with the topic to be delivered via the stream
define type mailbox
  topic
  deliveries stream of []byte

define context mailboxes as []mailbox

define function deliver chunk
  @msg = @chunk unwrap for context recipient

```

```

mailboxes each @mailbox                                13
    @payload = unseal @msg with @mailbox topic          14
    if @payload then                                    15
        write @msg payload                              16
        to @mailbox deliveries                          17

```

Definition 97 – pss API: send.

```

// /pss                                                    1
                                                                2
define function send @payload []byte                      3
    about topic                                           4
    for recipient                                         5
    ?to targets                                           6
has api POST on "/pss/<recipient>/<topic>(?targets=<targets>)" 7
    with @payload as body                                8
as                                                        9
    targets = lookup.targets for @recipient if no @targets 10
    context tag = tag/tag{}                               11
    seal @payload with @topic                             12
        wrap for @recipient                             13
            recipient
            to @targets                                  14
            trojan chunk
            store                                        15
        return tag                                     16
        // seal with topic
        // encrypt if given
        // mine nonce and returns
        // to be sent by push-sync
        // tag to monitor status

```

Definition 98 – pss API: receive.

```

// /pss                                                    1
                                                                2
define function receive about topic                      3
    on uint64 @channel                                    4
has api POST on "/pss/subscribe/<topic>(?on=<channel>)" 5
as                                                        6
    @stream = open @channel                              7
    context mailboxes append= mailbox{ @topic, @stream } 8
                                                                9
define function cancel topic                            10
    on @channel uint64                                   11
has api DELETE on "/pss/subscribe/<topic>(?on=<channel>)" 12
as                                                        13

```



```
context mailboxes any ch
```

14

2.7.3 Envelopes

Definition 99 – Envelope.

```
// /pss  
  
define type envelope  
  id [segment size]byte  
  sig crypto/signature  
  ps postage/stamp
```

1

2

3

4

5

6

2.7.4 Update notifications

2.7.5 Chunk recovery

Definition 100 – Targeted delivery request.

```
// /targeted.delivery  
  
define type envelope  
  id segment  
  signature crypto/signature  
  
define function wrap address  
  ?by @key crypto  
  to @targets []target  
as  
  @key = crypto/random if no @key  
  @n = crypto/mine @n such that  
    @targets any is prefix of  
      hash @n and @key account  
  @sig = crypto/sign hash @n and @address  
    by @key  
  envelope{ @n, @sig }
```

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

17

18

Definition 101 – Prod missing chunk notification with recovery request.

c

```
// /recovery 1
2
define type request 3
    address 4
    envelope 5
6
define function request address 7
    with @response bool 8
    to @targets 9
as 10
    @request = request{ @address } 11
    if @response then 12
        @request envelope = 13
            targeted.delivery/wrap @address by @key to @targets 14
    pss/send @request bytes about "RECOVERY" to @targets 15
```

Part I

Appendix

Appendix A

Density-based size estimation

Nodes in Swarm must utilise their full reserve capacity: nodes will potentially further replicate chunks in case they have unused reserve capacity beyond storing their share necessary for a system-wide level of redundancy required. To incentivise this, participating in the redistribution game involves a check called *proof of resources* which is supposed to verify the size of reserve from which the reserve samples are generated. The insight here is that the sample is the lowest range of a uniformly random variate over the entire 256-bit address space. Intuitively, the higher the original volume of the sampled set, the denser it is, the lower the expected maximum value in the sample. Conversely, a constraint on the maximum value of the last element in the sample practically puts a minimum cardinality requirement on the sampled set using a solution called *density based set size estimation*.

We are given n independent, uniformly distributed values between 0 and 1.¹ Let the value of the k th smallest of these be x_k (so the smallest of the n values is x_1 , the second smallest is x_2 , and so on, up to x_n). What is the distribution of x_k , given n ? And what is the threshold value u such that for any given probability α , the chance of obtaining an x_k lower than u is α ?

Note that the distance between any two adjacent values out of n independent uniform variates follows an exponential distribution, as long as n is sufficiently large.² The rate parameter of this exponential distribution is $n + 1$, where n is increased by one to account for the fact

¹Because we work with densities, the actual integer range is not relevant and results obtained for the unit interval can simply be rescaled to the Swarm use case by multiplying with 2^{256} .

²This follows from the fact that n independent uniform variates can be thought of as realizing a Poisson process, whereby the timing of events is random, and it is known that the nearest-neighbour distribution (i.e., waiting time between two consecutive events) is then exponentially distributed.

that the expected gap between adjacent values is $1/(n+1)$.³

We are after the distribution of the k th value,⁴ x_k , which then can be thought of as arising from the sum of k independent exponential variables, each with a rate parameter of $n+1$. This is known to result in an Erlang distribution with shape parameter k and rate parameter $n+1$. Using $X(\lambda)$ to denote the exponential distribution with rate λ and $E(k, \lambda)$ to denote the Erlang distribution with shape k and rate λ :

$$\sum_{i=1}^k X_i(n+1) = E(k, n+1), \quad (\text{A.1})$$

where the subscript i in $X_i(n+1)$ distinguishes between independent exponentially distributed random variables. The probability density function $E(x, k, n+1)$ of the Erlang distribution itself is given by

$$E(x, k, n+1) = \frac{(n+1)^k x^{k-1} e^{-(n+1)x}}{(k-1)!}. \quad (\text{A.2})$$

This distribution contains the answer to the first question: what is the distribution of the k th smallest value out of n independent uniform variates between 0 and 1? For example, if $k=16$ and n is either 500, 750, or 1000, we get the distributions shown in Figure A.1.

We can now answer the second question: given n and a confidence level α , what is the threshold value u for x_k such that the probability that $x_k < u$ is equal to α ? That is, we wish to know the value $x = u$ at which the probability distribution has encompassed a given area of α (see figure A.2).

The area under the curve of the Erlang distribution is given by its cumulative distribution function $P(x, k, n+1)$, which is known to be

$$P(x, k, n+1) = \int_0^x E(y, k, n+1) dy = \frac{1}{(k-1)!} \int_0^{(n+1)x} t^{k-1} e^{-t} dt. \quad (\text{A.3})$$

The latter expression is sometimes written as $\tilde{\gamma}(k, (n+1)x)$, where

$$\tilde{\gamma}(k, x) = \frac{1}{\Gamma(k)} \int_0^x t^{k-1} e^{-t} dt \quad (\text{A.4})$$

³For $n=2$, the mean outcome is $x_1 = 1/3$ and $x_2 = 2/3$; for $n=3$, it is $x_1 = 1/4$, $x_2 = 2/4$, $x_3 = 3/4$; and so on: for arbitrary n , $x_i = i/(n+1)$, with the gap between adjacent values in this ideal case always being $1/(n+1)$.

⁴An alternative approach using order statistic expresses x_k via a beta distribution. It is very difficult to prove that the Beta distribution's quantile function is a strictly decreasing function of n , which is a key piece of the argument presented here. Although this method is exact even for small n , in our case, n always a very large number, therefore we adopted the other method.




Figure A.1: Probability density function $E(x, k, n + 1)$ of the Erlang distribution, with $k = 16$ and n either 500, 750, or 1000.




Figure A.2: The distribution of x_{16} , i.e., the 16th smallest value from among $n = 500$ independently and uniformly drawn variates between 0 and 1. The area under the curve is shaded up to 5% of its area. The point at which the shading stops is therefore the value u for which there is only a 5% chance of getting an even smaller x_{16} .

is the regularized lower incomplete gamma function. We therefore want to solve the equation

$$\alpha = \int_0^u E(y, k, n + 1) dy = P(u, k, n + 1) \quad (\text{A.5})$$

for u .

Inverting this expression in u (since the cumulative distribution function increases monotonically in u , the inverse exists) leads to the quantile function $Q(\alpha, k, n + 1)$ of the Erlang distribution: $u = Q(\alpha, k, n + 1)$. The quantile function is known to be expressible as

$$Q(\alpha, k, n + 1) = \frac{\tilde{\gamma}^{-1}(k, x)}{n + 1}, \quad (\text{A.6})$$

where $\tilde{\gamma}^{-1}(k, x)$ is the inverse regularized lower incomplete gamma function. Its particular form is of no interest to us, except for two properties. First, it is positive for all x .⁵ Second, it is independent of n . Instead, the entire dependence of $Q(\alpha, k, n + 1)$ on n is given by the $n + 1$ term in the denominator of Equation A.6. From this, we conclude that $Q(\alpha, k, n + 1)$ is a strictly decreasing function of n .

These two points lead to an important consequence. Say we compute the threshold u for a given α and n in order to have an upper bound on a lower quantile. Now, if we were to decrease n but hold all other things equal, the threshold will always get higher than what it was before. The threshold obtained for higher values of n may therefore serve as a conservative estimate of the threshold for lower values: if u is a threshold such that the k th smallest out of n uniform variates is only smaller than u in α of cases, then for any amount $m < n$, the chance of the k th variate conforming to the same constraint (i.e., $x_k < u$) is now even smaller than α .

Conversely, if we were to constrain x_k so that the probability of not getting a value smaller than u is lower than β (minimising a higher quantile), we find that the constraint remains true as n is increased.

Armed with these results, let us see how Equation A.6 can be used for the estimation procedure. There are two problems to tackle, ultimately relating to the two aspects of a test's accuracy. First, we want to catch inadequate storers slacking on volume. In other words, we want to constrain the x_k values so that we can safely say that any attacker with a stored volume below an acceptable size n has a probability less than α to obtain such a small x_k by pure chance. Construing the condition for $x_k < u$ as a test to filter honest players (just based on the size of their reserve), $1 - \alpha$ expresses the *sensitivity* of the test. From the previous argument on the monotonic dependence of α on n , it is safe to use a condition that requires x_k to stay below a threshold obtained for n .

Second, we want to avoid situations when honest participants end up not satisfying the above constraint even though they sampled from a set larger than the required minimum. Given a target volume $m > n$, the error rate of false negatives is guaranteed to be less than

⁵This stands to reason: the quantile function of a distribution on $x \in [0, \infty)$ is itself between 0 and ∞ , and $\tilde{\gamma}^{-1}(k, x)$ is just the quantile function of the Erlang distribution times the positive constant $n + 1$.

β obtained from the quantile function with parameters m, k , and u . The quantity $1 - \beta$ is the *specificity* or *precision* of the test.

Figure A.3 illustrates this idea, for two different distributions in both the lower- and upper-end estimation. What we want is to choose u to simultaneously make sure that dishonest players do not sneak through the system *and* also that honest players do not get excluded too often. This translates to make both α and β as small as possible.




Figure A.3: Recall and precision of proof of reserve size validation: Any chosen u will lead to different α and β values, depending on n . Here u is fixed at 0.02. The top panel shows distributions for $x \equiv x_{16}$ with $n = 500$ (blue) and $n = 1000$ (yellow). The area left under the blue curve to the left of u is equal to α (blue shade); the area under the yellow curve right of u is equal to β (yellow shade). If the curves overlap considerably (top), it is impossible to choose an u such that α and β are simultaneously small.

One way to try and find the best compromise is by minimising $\alpha + \beta$ (the *accuracy* of the test) and pick the u value at the optimum to be used in the proof of resources test. To this end, one can vary α between 0 and 1 and, for each of its values, solve the equation $\alpha = Q(1 - \beta, k, n + 1)$ (where n is the larger value, used for estimating β). This way, we get a β value for every possible α . Then, we can find the combination which minimises $\alpha + \beta$, and determine the value of u that leads to this optimum. As illustrated in figure A.4, larger values of k yield a trade-off curve along which better accuracies can be achieved.




Figure A.4: Optimal accuracy of reserve size probe By increasing k , one can get better optima for minimising $\alpha + \beta$. Here $n = 10^6$ for estimating α and $2 \cdot 10^6$ for estimating β , and k is either 8, 16, or 32 (colors). The values of u associated with the optima are also shown.

Table A.1 summarizes the important numerical results in Figure A.4.⁶

k	α	β	u	$u \cdot 2^{256}$
8	0.196216	0.1373622	$5.54589 \cdot 10^{-6}$	$6.421705 \cdot 10^{71}$
16	0.097612	0.0716570	$1.10923 \cdot 10^{-5}$	$1.284401 \cdot 10^{72}$
32	0.029386	0.0219151	$2.21962 \cdot 10^{-5}$	$2.570140 \cdot 10^{72}$

Table A.1: Proof of density parameter calibration. Assuming $n = 10^6$ and $m = 2 \cdot 10^6$ to calculate recall and precision error rates α and β , respectively, the cutoff value for the proof is calibrated by optimizing on accuracy using sample sizes 8, 16, 32.

⁶Since the hash function used to generate random variates does so in the range $[0, 2^{256} - 1]$ instead of $[0, 1]$, the calculated thresholds are scaled with 2^{256} to show where they would fall in their actual range.

Appendix B

Source of randomness

As the *neighbourhood selection anchor* will directly affect which neighbourhood wins the pot, it is prudent to derive the randomness from a source of entropy that cannot be manipulated. A common solution to obtain randomness is to have independent parties committing to a random nonce with a stake. The random seed for a round is defined as the xor of all revealed nonces. Given the nonces are independent sources fixed in the commit, no individual participant has the ability to skew randomness by selecting a particular nonce. Thanks to the commutativity of xor, the order of reveals is also irrelevant. However, if the reveal transactions are sequential, committers compete at holding out since the last one to reveal can effectively choose the resulting seed to be either including its committed nonce (if they do reveal) or not (if they do not reveal). The threat to slash the stake of non-revealers serves to eliminate this degree of freedom from last revealers and thus renders this scheme a secure random oracle assuming there is at least one honest (non-colluding) party.¹

Now note that the redistribution scheme already has a commit reveal scheme as well as stake slashed for non-revealers, so a potential random oracle is already part of the proposed scheme. Incidentally, the beginning of the claim phase is when new randomness is needed to select the truth and a winner. Importantly, these random values are only needed if there is a claim which implies that there were some commits and reveals to choose from. Or conversely, if there are no reveals in the round,² the random seed is undefined but is also not needed for the claim.

The random seed that transpires at the beginning of the claim phase can serve as the *reserve*

¹If the stake is higher than the reward pot, one cannot afford being slashed with even just one commit without a loss. If this cannot be guaranteed, slashing of the stake is not an effective deterrent.

²If saboteurs get slashed or frozen in the claim transaction, if there is no claim, the committers get away without being punished. This can be remedied if the staking contract keeps a flag on each overlay (set when commits, unsets when reveals in the same round) and the check and punishment happens as a result of a commit call in the case the flag is found set.

sample salt (nonce input to modified hash used in the sampling) with which the nodes in the selected neighbourhood can start calculating their reserve sample.

The neighbourhood for the next round is selected by the *neighbourhood selection anchor*, which is, similarly to the truth and winner selection nonces, deterministically derived from the same random seed. Unlike the nonces used to select from the reveals, neighbourhood selection should be well defined for the following round even if a round is skipped, i.e., when there is no reveals. To cover this case skipped rounds keep the random seed of the previous round. However, in order to rotate selected neighbourhoods through skipped rounds, we derive the neighbourhood selection anchor from the seed by factoring in the number of game rounds passed since the last reveal.³

In order to provide protection against the case when each committer in the neighbourhood is colluding, and can afford losing stake we need to make sure that the entropy is still high otherwise the nodes can influence the neighbourhood selection nonce and reselect themselves or a fixed colluding neighbourhood (or increase the chances of reselection).

Definition 103 – Random seed for the round.

Define the random seed of the round as the xor of all obfuscation keys sent as part of the reveal transaction data during the entire reveal period:

$$\mathcal{R} \quad : \quad \Gamma \mapsto \text{Nonce} \tag{B.1}$$

$$\mathcal{R}(\gamma) \stackrel{\text{def}}{=} \begin{cases} \mathcal{R}(\text{Prev}(\gamma)) & \text{if } |\text{Reveals}(\gamma)| = 0 \\ \bigvee_{r \in \text{Reveals}(\gamma)} \text{NONCE}(r) & \text{otherwise} \end{cases} \tag{B.2}$$

Lemma 104 – Round seed is a secure random oracle.

The nonce produced by xoring the revealed obfuscation keys is a correct source of entropy.

Proof. Assuming n independent parties committing, choosing any particular nonce will leave the outcome fully random.

³Otherwise a selected neighbourhood could collude maliciously not to commit/reveal and have the pot roll over to the following round. By simply holding out for a number of redistribution rounds, they could unfairly multiply their reward when they eventually claim the pot.

Appendix C

Parameter constants

Table C.1 lists the constants used by Swarm with their type, default value and description.

CONSTANT NAME	TYPE	VALUE	DESCRIPTION
BZZ_NETWORK_ID	<i>uint64</i>	0	ID of the swarm network
PHASE_LENGTH	<i>uint64</i>	38	length of commit phase of the redistribution game round in number of blocks
ROUND_LENGTH	<i>uint64</i>	152	length of one redistribution game round in number of blocks, fixed at 4 times the PHASE_LENGTH
NODE_RESERVE_DEPTH	<i>uint8</i>	23	size requirement for client reserve capacity given in base 2 log number of chunks
SAMPLE_DEPTH	<i>uint8</i>	4	base 2 log number of chunks in sample
MAX_SAMPLE_VALUE	<i>uint256</i>	$1.2844 \cdot 10^{72}$	maximum value for last transformed address in reserve sample, i.e., < 1% chance the sampled set size is below a fourth of the prescribed node reserve size.
MINIMUM_STAKE	<i>uint256</i>	10	minimum stake amount in BZZ to be re-defined as minimum stake given in storage rent units
MIN_STAKE_AGE	<i>uint256</i>	228	minimum number of blocks stakers need to wait after update or creation for the stake to be useable. Defaults to one and a half rounds to prevent opportunistic manipulation of stake after a neighbourhood is selected.

PRICE_2X_ROUNDS	<i>uint64</i>	64	number of rounds it takes for the price to double in the presence of a consistent lowest degree signal of undersupply.
NHOOD_PEER_COUNT	<i>uint8</i>	4	minimum number of nodes required to form a fully connected neighbourhood.

Table C.1: Parameter constants